

CS370 Operating Systems

Colorado State University

Yashwant K Malaiya

Fall 2016 Lecture 7



Slides based on

- Text by Silberschatz, Galvin, Gagne
- Various sources

CS370 OS Ch3 Processes

- Process Concept: a program in execution
- Process Scheduling
- Processes creation and termination
- Interprocess Communication using shared memory and message passing

FAQ

- Program control Block – What does it include and why? Where is it saved? another look
- Process Scheduling – Why? How? more soon
- How do multiple processes, that are part of a single program, interact. Inter-process communication details coming up.
- What happens to a process once it is finished? Resources deallocated, but only after ..
- If a process forks a child, what happens to the parent? It continues.
- When the CPU is running user-processes how does the OS run? Is the kernel one process? Where does it reside?
- Can a process induce its own context switch? Yes, we'll see how.


Notes

- Please register your iClicker: Canvas > iClicker menu item
 - Bring to class
- Introduce yourself on Canvas
- Canvas Discussion board available for discussions
 - usual rules apply
- The TA office hours and photos are on cs270 home page

More FAQ

- Q3.1 Registers are managed by the (a) Compiler
- Q2.b: Multiprogramming requires presence of multiple processors.
False
- Q4.b POSIX API are used in a high level language. True

Diagram of Process State



Ready to Running: scheduled by scheduler

Running to Ready: scheduler picks another process, back in ready queue


Running to Waiting (Blocked) : process blocks for input/output

Waiting to Ready: Input available

Process Control Block (PCB)


Information associated with each process
(also called **task control block**)

- Process state – running, waiting, etc
- Program counter – location of instruction to next execute
- CPU registers – contents of all process-centric registers
- CPU scheduling information- priorities, scheduling queue pointers
- Memory-management information – memory allocated to the process
- Accounting information – CPU used, clock time elapsed since start, time limits
- I/O status information – I/O devices allocated to process, list of open files



Representation of Process Scheduling

- **Queueing diagram** represents queues, resources, flows



Schedulers

- **Short-term scheduler** (or **CPU scheduler**) – selects which process should be executed next and allocates CPU
 - Sometimes the only scheduler in a system
 - Short-term scheduler is invoked frequently (10-100 milliseconds) \Rightarrow (must be fast)
- **Long-term scheduler** (or **job scheduler**) – selects which processes should be brought into the ready queue. Batch systems.
 - invoked infrequently (seconds, minutes) \Rightarrow (may be slow)
 - controls the **degree of multiprogramming**
- Processes can be described as either:
 - **I/O-bound process** – spends more time doing I/O than computations, many short CPU bursts
 - **CPU-bound process** – spends more time doing computations; few very long CPU bursts
- Long-term scheduler strives for good *process mix*

Scheduling

- **UNIX and Windows systems often have no long-term scheduler**
- **Medium-term scheduler** can be added if degree of multiple programming needs to decrease
 - Remove process from memory, store on disk, bring back in from disk to continue execution: **swapping**
 - Handles strains on memory etc.

Multitasking in Mobile Systems

- Some mobile systems (e.g., early version of iOS) allow only one process to run, others suspended
- Due to screen real estate, user interface limits iOS provides for a
 - Single **foreground** process- controlled via user interface
 - Multiple **background** processes– in memory, running, but not on the display, and with limits
 - Limits include single, short task, receiving notification of events, specific long-running tasks like audio playback
 - iPad with IOS9: Split View multitasking
- Android runs foreground and background, with fewer limits
 - Background process uses a **service** to perform tasks
 - Service can keep running even if background process is suspended
 - Service has no user interface, small memory use



Context Switch


- When CPU switches to another process, the system must **save the state** of the old process and load the **saved state** for the new process via a **context switch**
- **Context** of a process represented in the PCB
 - CPU registers, process state, memory management info etc.
- Context-switch time (typically microsecs) is overhead; the system does no useful work while switching
 - The more complex the OS and the PCB → the longer the context switch
- Time dependent on hardware support
 - Some hardware provides multiple sets of registers per CPU → multiple contexts loaded at once

Processes creation & termination

Process Creation

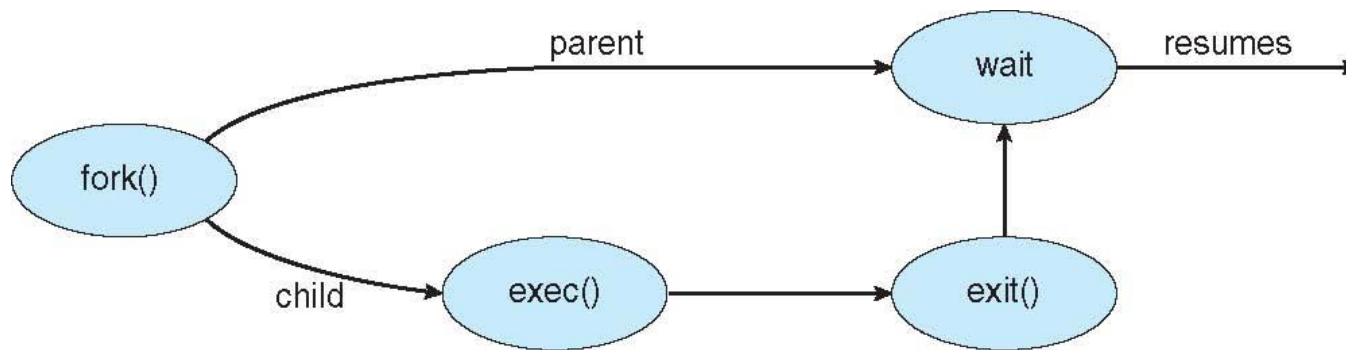
- **Parent** process create **children** processes, which, in turn create other processes, forming a **tree** of processes
- Generally, process identified and managed via a **process identifier (pid)**
- Resource sharing options
 - Parent and children share all resources
 - Children share subset of parent's resources
 - Parent and child share no resources
- Execution options
 - Parent and children execute concurrently
 - Parent waits until children terminate

A Tree of Processes in Linux



Process Creation (Cont.)

- Address space
 - Child duplicate of parent
 - Child has a program loaded into it
- UNIX examples
 - **fork()** system call creates new process
 - **exec()** system call used after a **fork()** to replace the process' memory space with a new program



Fork () to create a child process

- Fork creates a copy of process
- Return value from fork (): integer
 - When > 0 :
 - Running in (original) **Parent** process
 - return value is pid of new child
 - When $= 0$:
 - Running in new **Child** process
 - When < 0 :
 - Error! Perhaps exceeds resource constraints. sets errno (a global variable in errno.h)
 - Running in original process
- All of the state of original process duplicated in both Parent and Child!
 - Memory, File Descriptors (next topic), etc...

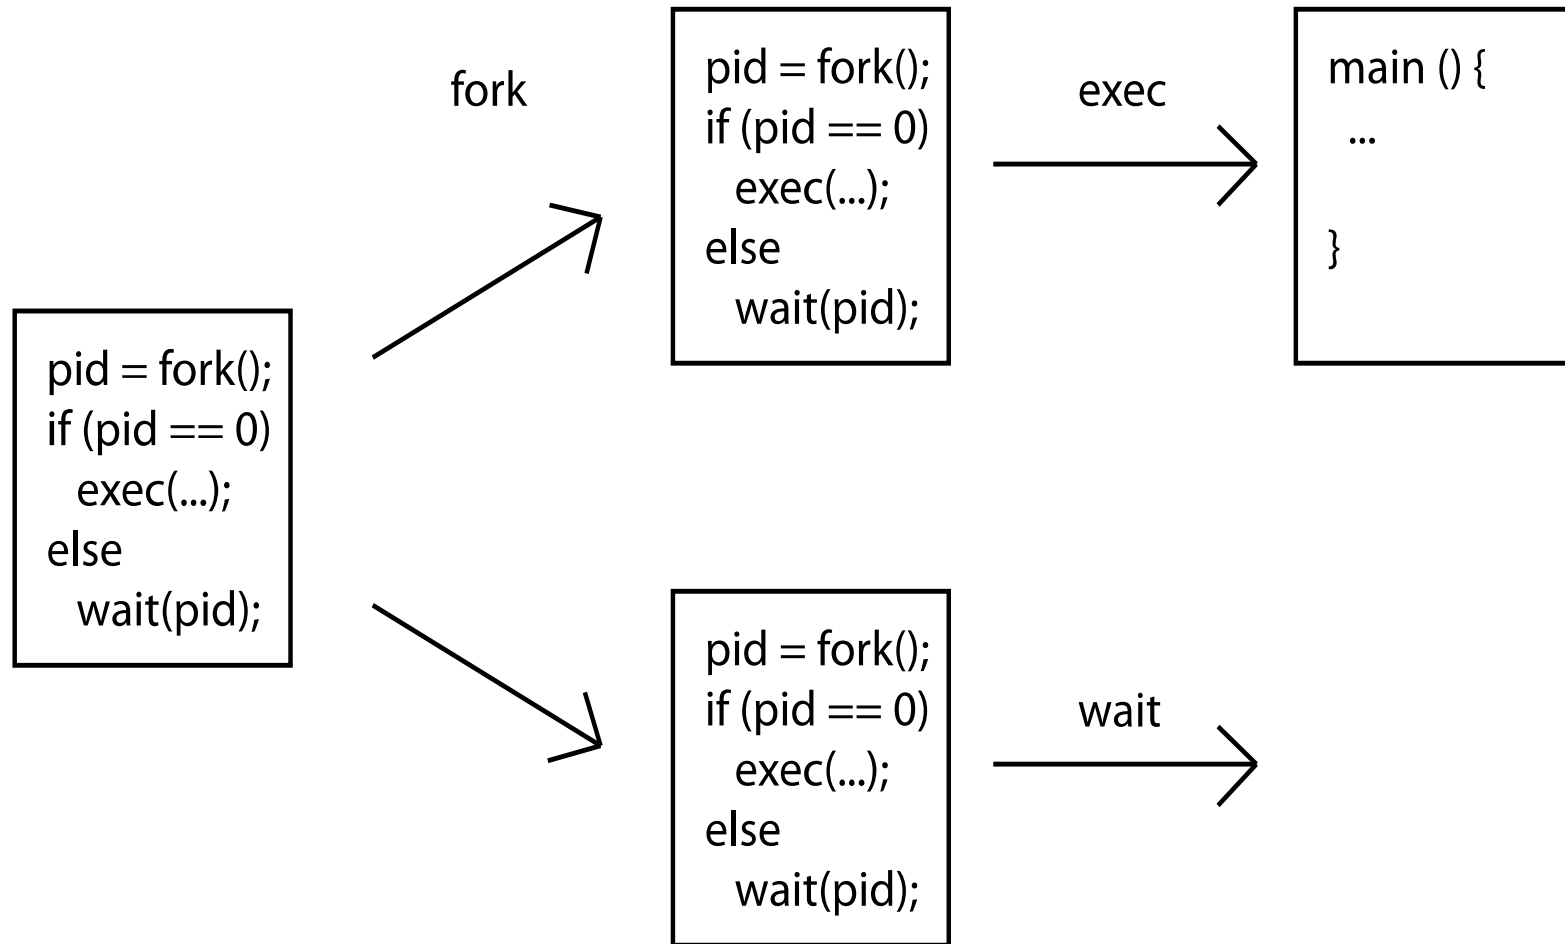
Process Management System Calls

- UNIX fork – system call to create a copy of the current process, and start it running
 - No arguments!
- UNIX exec – system call to *change the program* being run by the current process. Several variations.
- UNIX wait – system call to wait for a process to finish
- Details: see man pages

Notes:

```
pid_t pid = getpid(); /* get current processes PID */;
waitpid(cid, 0, 0); /* Wait for my child to terminate. */
exit (0); /* Quit*/
kill(cid, SIGKILL); /* Kill child*/
```

UNIX Process Management



C Program Forking Separate Process

```
#include <sys/types.h>
#include <stdio.h>
#include <unistd.h>

int main()
{
    pid_t pid;

    /* fork a child process */
    pid = fork();

    if (pid < 0) { /* error occurred */
        fprintf(stderr, "Fork Failed");
        return 1;
    }
    else if (pid == 0) { /* child process */
        execlp("/bin/ls", "ls", NULL);
    }
    else { /* parent process */
        /* parent will wait for the child to complete */
        wait(NULL);
        printf("Child Complete");
    }

    return 0;
}
```

execlp(3) - Linux man page
<http://linux.die.net/man/3/execlp>

Forking PIDs

```
#int main()
{
    pid_t cid;

    /* fork a child process */
    cid = fork();
    if (cid < 0) { /* error occurred */
        fprintf(stderr, "Fork Failed\n");
        return 1;
    }
    else if (cid == 0) { /* child process */
        printf("I am the child %d, my PID is %d\n", cid, getpid());
        execlp("/bin/ls", "ls", NULL);
    }
    else { /* parent process */
        /* parent will wait for the child to complete */
        printf("I am the parent with PID %d, my parent is %d, my child is %d\n", getpid(), getppid(), cid);
        wait(NULL);

        printf("Child Complete\n");
    }

    return 0;
}
```

```
Ys-MacBook-Air:ch3 ymalaiya$ ./newproc-posix_m
I am the parent with PID 494, my parent is 485, my child is 496
I am the child 0, my PID is 496
DateClient.java                               newproc-posix_m

Child Complete
Ys-MacBook-Air:ch3 ymalaiya$
```

wait/waitpid

- Wait/waitpid () allows caller to suspend execution until child's status is available
- Process status availability
 - Generally after termination
 - Or if process is stopped
- `pid_t waitpid(pid_t pid, int *status, int options);`
- The value of pid can be:
 - 0 wait for any child process with same *process group ID* (perhaps inherited)
 - > 0 wait for child whose process group ID is equal to the value of pid
 - Others
- Status: where status info needs to be saved

Process Termination

- Process executes last statement and then asks the operating system to delete it using the **exit()** system call.
 - Returns status data from child to parent (via **wait()**)
 - Process' resources are deallocated by operating system
- Parent may terminate the execution of children processes using the **abort()** system call. Some reasons for doing so:
 - Child has exceeded allocated resources
 - Task assigned to child is no longer required
 - The parent is exiting and the operating systems does not allow a child to continue if its parent terminates

Process Termination

- Some operating systems do not allow child to exist if its parent has terminated. If a process terminates, then all its children must also be terminated.
 - **cascading termination.** All children, grandchildren, etc. are terminated.
 - The termination is initiated by the operating system.
- The parent process may wait for termination of a child process by using the **wait()** system call. The call returns status information and the pid of the terminated process

```
pid = wait(&status);
```
- If no parent waiting (did not invoke **wait()**) process is a **zombie**
- If parent terminated without invoking **wait**, process is an **orphan**

Multiprocess Architecture – Chrome Browser

- Early web browsers ran as single process
 - If one web site causes trouble, entire browser can hang or crash
- Google Chrome Browser is multiprocess with 3 different types of processes:
 - **Browser** process manages user interface, disk and network I/O
 - **Renderer** process renders web pages, deals with HTML, Javascript. A new renderer created for each website opened
 - Runs in **sandbox** restricting disk and network I/O, minimizing effect of security exploits
 - **Plug-in** process for each type of plug-in

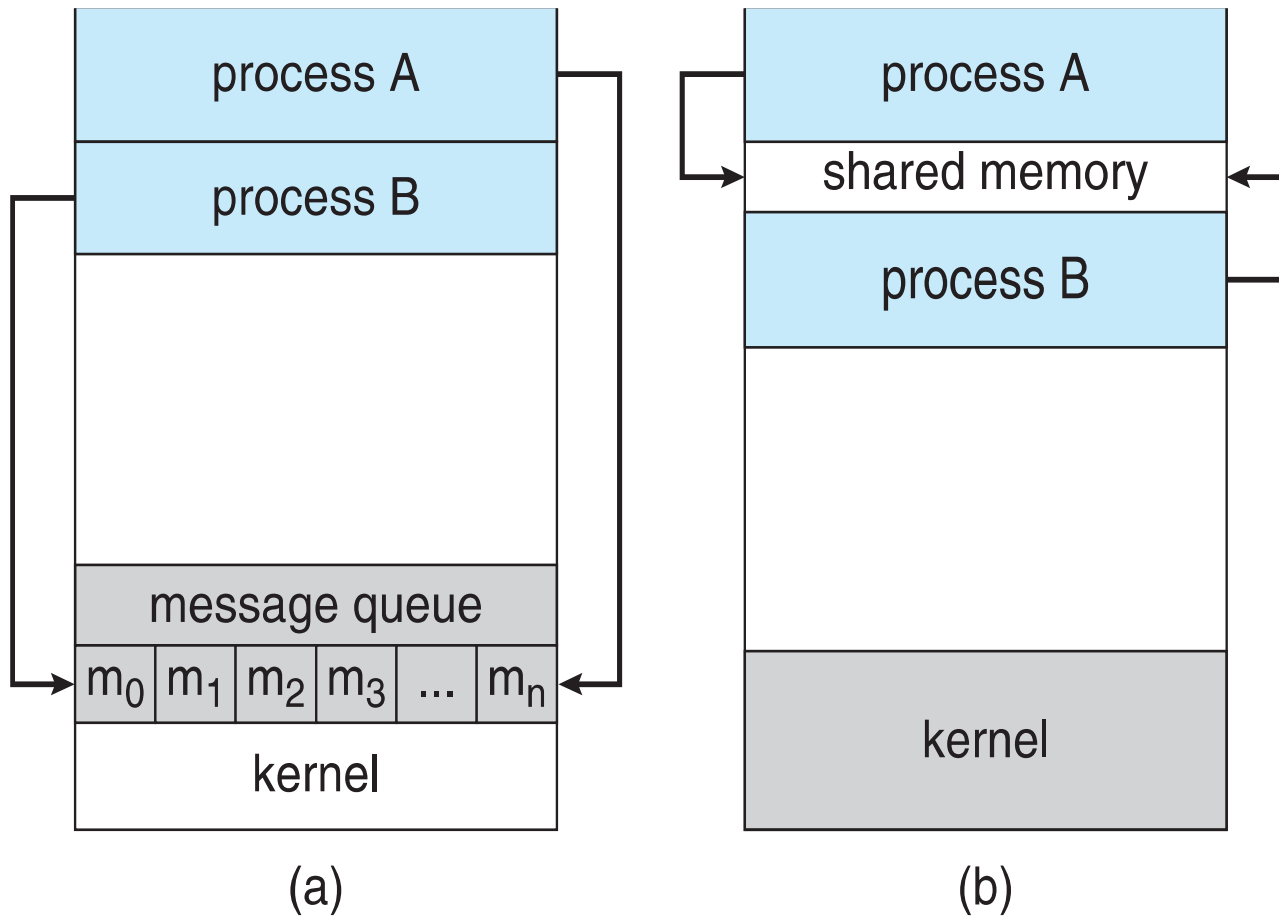


Interprocess Communication

- Processes within a system may be *independent* or *cooperating*
- Cooperating process can affect or be affected by other processes, including sharing data
- Reasons for cooperating processes:
 - Information sharing
 - Computation speedup
 - Modularity
 - Convenience
- Cooperating processes need **interprocess communication (IPC)**
- Two models of IPC
 - **Shared memory**
 - **Message passing**

Communications Models

(a) Message passing. (b) shared memory.



Cooperating Processes

- ***Independent*** process cannot affect or be affected by the execution of another process
- ***Cooperating*** process can affect or be affected by the execution of another process
- Advantages of process cooperation
 - Information sharing
 - Computation speed-up
 - Modularity
 - Convenience