

CS370 Operating Systems

Colorado State University

Yashwant K Malaiya

Fall 2016 Lecture 9



Slides based on

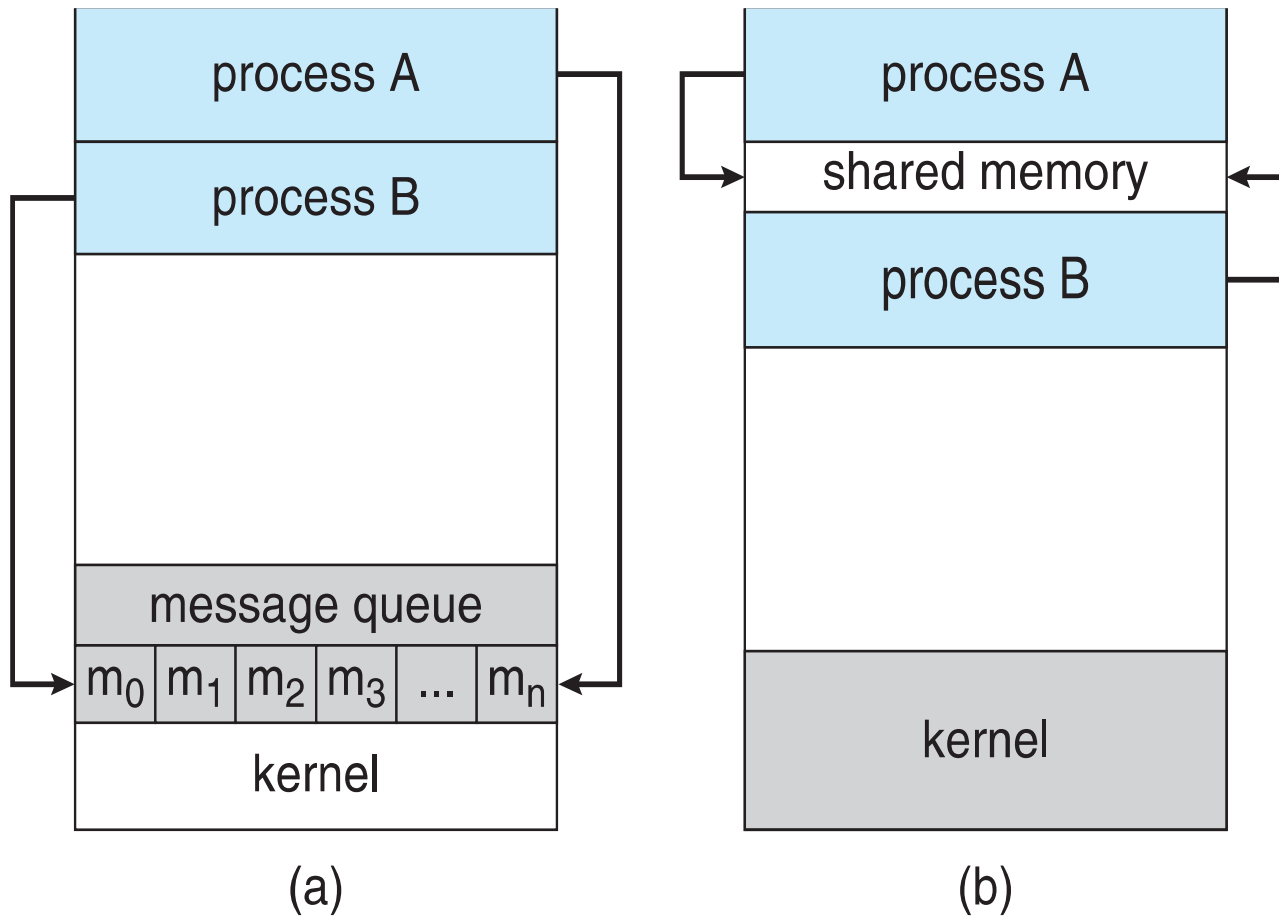
- Text by Silberschatz, Galvin, Gagne
- Various sources

FAQ

- Why would users want to use `fork()`, `wait()` etc in your own code, since OS does a good job?
- Is the child [initially] an exact copy of the parent? Does it have the same values in heap, stack, and program counter?
- **Zombie:** process has exited, but parent hasn't called `wait()` to obtain exit status. Two bytes.
- **Process Group ID:** a process can send a signal to all members of the PGID.
- Are unbounded buffers used in any systems?
- Why examples use producers and consumers in an infinite loop?


Communications Models

(a) Message passing. (b) shared memory.



Interprocess Communication – Shared Memory

- An area of memory shared among the processes that wish to communicate
- The communication is under the control of the users processes not the operating system.
- Major issues is to provide mechanism that will allow the user processes to **synchronize** their actions when they access shared memory.
 - Synchronization is discussed in great details in Chapter 5.
- Example soon.



Only one process
may access
shared memory
at a time

- Mechanism for processes to communicate and to synchronize their actions
- Message system – processes communicate with each other without resorting to shared variables
- IPC facility provides two operations:
 - **send**(*message*)
 - **receive**(*message*)
- The *message* size is either fixed or variable

Message Passing (Cont.)

- If processes P and Q wish to communicate, they need to:
 - Establish a **communication link** between them
 - Exchange messages via send/receive
- Implementation issues:
 - How are links established?
 - Can a link be associated with more than two processes?
 - How many links can there be between every pair of communicating processes?
 - What is the capacity of a link?
 - Is the size of a message that the link can accommodate fixed or variable?
 - Is a link unidirectional or bi-directional?

- Implementation of communication link
 - Physical:
 - Shared memory
 - Hardware bus
 - Network
 - Logical: Options (details next)
 - Direct (process to process) or indirect (mail box)
 - Synchronous (blocking) or asynchronous (non-blocking)
 - Automatic or explicit buffering

Direct Communication

- Processes must name each other explicitly:
 - **send** (*P, message*) – send a message to process P
 - **receive**(*Q, message*) – receive a message from process Q
- Properties of communication link
 - Links are established automatically
 - A link is associated with exactly one pair of communicating processes
 - Between each pair there exists exactly one link
 - The link may be unidirectional, but is usually bi-directional

Indirect Communication

- Messages are directed and received from mailboxes (also referred to as ports)
 - Each mailbox has a unique id
 - Processes can communicate only if they share a mailbox
 - May be owned by a process (to receive) or the OS
- Properties of communication link
 - Link established only if processes share a common mailbox
 - A link may be associated with many processes
 - Each pair of processes may share several communication links
 - Link may be unidirectional or bi-directional

Indirect Communication

- Operations
 - create a new mailbox (port)
 - send and receive messages through mailbox
 - destroy a mailbox
- Primitives are defined as:
 - send**(*A, message*) – send a message to mailbox A
 - receive**(*A, message*) – receive a message from mailbox A

Indirect Communication

- Mailbox sharing
 - P_1 , P_2 , and P_3 share mailbox A
 - P_1 , sends; P_2 and P_3 receive
 - Who gets the message?
- Possible Solutions
 - a link to be associated with at most two processes
 - only one process at a time to execute a receive operation
 - system to select arbitrarily the receiver.
(Sender is notified who the receiver was.)

Synchronization(*blocking or not*)

- Message passing may be either blocking or non-blocking
- **Blocking** is termed **synchronous**
 - **Blocking send** -- sender is blocked until message is received
 - **Blocking receive** -- receiver is blocked until a message is available
- **Non-blocking** is termed **asynchronous**
 - **Non-blocking send** -- sender sends message and continues
 - **Non-blocking receive** -- the receiver receives:
 - A valid message, or
 - Null message
- Different combinations possible
 - If both send and receive are blocking, we have a **rendezvous**.
 - Producer-Consumer problem: Easy both block

Buffering

- Queue of messages attached to the link.
- implemented in one of three ways
 1. Zero capacity – no messages are queued on a link.
Sender must wait for receiver (rendezvous)
 2. Bounded capacity – finite length of n messages
Sender must wait if queue full
 3. Unbounded capacity – infinite length
Sender never waits

Examples of IPC Systems - POSIX

- Older scheme (System V) using `shmget()`, `shmat()`, `shmdt()`, `shmctl()`
- POSIX Shared Memory
 - Process first **creates shared memory segment**
`shm_fd = shm_open(name, O_CREAT | O_RDWR, 0666);`
 - Returns file descriptor (int) which identifies the file
 - Also used to open an existing segment to share it
 - Set the size of the object
`ftruncate(shm_fd, 4096);`
 - map the shared memory segment **in the address space of the process**
`ptr = mmap(0, SIZE, PROT_READ | PROT_WRITE, MAP_SHARED, shm_fd, 0);`
 - Now the process could write to the shared memory
`sprintf(ptr, "Writing to shared memory");`

Examples of IPC Systems - POSIX

■ POSIX Shared Memory

- Other process opens shared memory object `name`

```
shm_fd = shm_open(name, O_RDONLY, 0666);
```

- Returns file descriptor (int) which identifies the file

- map the shared memory object

```
ptr = mmap(0, SIZE, PROT_READ, MAP_SHARED, shm_fd, 0);
```

- Now the process can read from to the shared memory object

```
printf("%s", (char *)ptr);
```

- remove the shared memory object

```
shm_unlink(name);
```

IPC POSIX Producer

```
#include <stdio.h>
#include <stdlib.h>
#include <string.h>
#include <fcntl.h>
#include <sys/shm.h>
#include <sys/stat.h>

int main()
{
    /* the size (in bytes) of shared memory object */
    const int SIZE = 4096;
    /* name of the shared memory object */
    const char *name = "OS";
    /* strings written to shared memory */
    const char *message_0 = "Hello";
    const char *message_1 = "World!";

    /* shared memory file descriptor */
    int shm_fd;
    /* pointer to shared memory object */
    void *ptr;

    /* create the shared memory object */
    shm_fd = shm_open(name, O_CREAT | O_RDWR, 0666);

    /* configure the size of the shared memory object */
    ftruncate(shm_fd, SIZE);

    /* memory map the shared memory object */
    ptr = mmap(0, SIZE, PROT_WRITE, MAP_SHARED, shm_fd, 0);

    /* write to the shared memory object */
    sprintf(ptr,"%s",message_0);
    ptr += strlen(message_0);
    sprintf(ptr,"%s",message_1);
    ptr += strlen(message_1);

    return 0;
}
```


IPC POSIX Consumer

```
#include <stdio.h>
#include <stdlib.h>
#include <fcntl.h>
#include <sys/shm.h>
#include <sys/stat.h>

int main()
{
    /* the size (in bytes) of shared memory object */
    const int SIZE = 4096;
    /* name of the shared memory object */
    const char *name = "OS";
    /* shared memory file descriptor */
    int shm_fd;
    /* pointer to shared memory object */
    void *ptr;

    /* open the shared memory object */
    shm_fd = shm_open(name, O_RDONLY, 0666);

    /* memory map the shared memory object */
    ptr = mmap(0, SIZE, PROT_READ, MAP_SHARED, shm_fd, 0);

    /* read from the shared memory object */
    printf("%s", (char *)ptr);

    /* remove the shared memory object */
    shm_unlink(name);

    return 0;
}
```

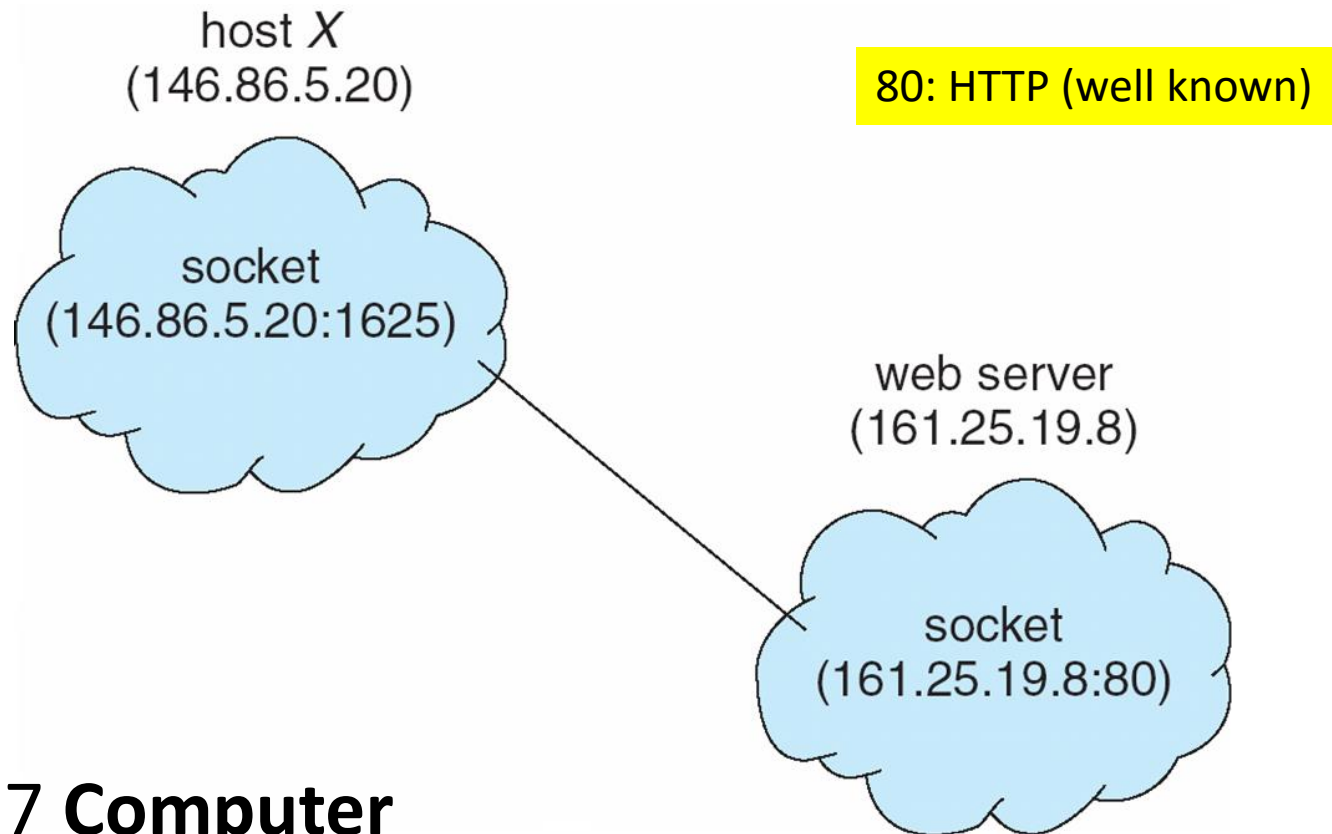
Communications in Client-Server Systems

- Sockets
- Remote Procedure Calls
- Pipes
- Remote Method Invocation (Java)

Sockets

- A **socket** is defined as an endpoint for communication
- Concatenation of IP address and **port** – a number included at start of message packet to differentiate network services on a host
- The socket **161.25.19.8:1625** refers to port **1625** on host **161.25.19.8**
- Communication consists between a pair of sockets
- All ports below 1024 are ***well known***, used for standard services
- Special IP address 127.0.0.1 (**loopback**) to refer to system on which process is running

Socket Communication



- **CS457 Computer Networks and the Internet**

Pipes

- Acts as a conduit allowing two processes to communicate
- One of the first IPC implementation mechanisms

Pipes


- Conduit allowing two processes to communicate
- Issues:
 - Is communication unidirectional or bidirectional?
 - If bidirectional, is it **half-duplex** (one way at a time) or **full-duplex** (both directions simultaneously)?
 - Must there exist a relationship (i.e., ***parent-child***) between the communicating processes?
 - Can the pipes be used over a network?

Pipes

- Command line:
 - Set up pipe between commands
ls | more
 - Output of ls delivered as input to more
- **Ordinary (“anonymous”) pipes** –Typically, a parent process creates a pipe and uses it to communicate with a child process that it created. Cannot be accessed from outside the process that created it.
- **Named pipes (“FIFO”)** – can be accessed without a parent-child relationship.

Ordinary Pipes


- Ordinary Pipes allow communication in standard producer-consumer style
- Producer writes to one end (the **write-end** of the pipe)
- Consumer reads from the other end (the **read-end** of the pipe)
- Ordinary pipes are therefore unidirectional (half duplex)
- **Require parent-child relationship** between communicating processes
- **pipe (int fd[])** to create pipe, fd[0] is the read-end, fd[1] is the write-end



- Windows calls these **anonymous pipes**

Ordinary Pipes

- ☐ Pipe is a special type of file.
- ☐ Inherited by the child
- ☐ Must close unused portions of the the pipe



UNIX pipe example

```
#define READ_END  0
#define WRITE_END 1
```

```
int fd[2];
```

create the pipe:

```
if (pipe(fd) == -1) {
    fprintf(stderr, "Pipe failed");
    return 1;
}
```

fork a child process:


```
pid = fork();
```

parent process:


```
/* close the unused end of the pipe */
close(fd[READ_END]);
```

```
/* write to the pipe */
write(fd[WRITE_END], write_msg, strlen(write_msg)+1);
```

```
/* close the write end of the pipe */
close(fd[WRITE_END]);
```



UNIX pipe example



child process:

```
/* close the unused end of the pipe */  
close(fd[WRITE_END]);
```

```
/* read from the pipe */  
read(fd[READ_END], read_msg, BUFFER_SIZE);  
printf("child read %s\n", read_msg);
```

```
/* close the write end of the pipe */  
close(fd[READ_END]);
```

Named Pipes

- Named Pipes (termed FIFO) are more powerful than ordinary pipes
- Communication is bidirectional
- No parent-child relationship is necessary between the communicating processes
- Several processes can use the named pipe for communication
- Provided on both UNIX and Windows systems