

# CS370 Operating Systems

Colorado State University

Yashwant K Malaiya

Fall 2016 Lecture 18



## Slides based on

- Text by Silberschatz, Galvin, Gagne
- Various sources

# FAQ

- Why are critical sections important?
  - Correctness, data corruption
- Can't critical sections cause starvation?
  - Not if they satisfy ..
- Two processes do not share any resources, do they need critical sections?
- Why we didn't study critical sections before?
- Are critical sections for two interacting processes the same length?

# Critical Section

```
do {
```

```
    entry section
```

```
        critical section
```

```
    exit section
```

```
        remainder section
```

```
} while (true);
```

Request permission  
to enter

Housekeeping to let  
processes to enter  
other

# Mutex Locks

- Previous solutions are complicated and generally inaccessible to application programmers
- OS designers build software tools to solve critical section problem
- Simplest is **mutex** lock
- Protect a critical section by first `acquire()` a lock then `release()` the lock
  - Boolean variable indicating if lock is available or not
- Calls to `acquire()` and `release()` must be atomic
  - Usually implemented via hardware atomic instructions
- But this solution requires **busy waiting**
  - This lock therefore called a **spinlock**

# acquire() and release()

```
acquire() {  
    while (!available)  
        ; /* busy wait */  
}
```

```
release() {  
    available = true;  
}
```

## •Usage

```
do {  
    acquire lock  
    critical section  
    release lock  
    remainder section  
} while (true);
```

# acquire() and release()

- Use board

# Semaphores by Dijkstra

- Synchronization tool that provides more sophisticated ways (than Mutex locks) for process to synchronize their activities.
- Semaphore **S** – integer variable
- Can only be accessed via two **indivisible (atomic)** operations

– **wait()** and **signal()**

- Originally called **P()** and **V()** based on Dutch words

- Definition of the **wait()** operation

```
wait(S) {  
    while (S <= 0)  
        ; // busy wait  
    S--;  
}
```

Waits until  
another process  
makes S=1

- Definition of the **signal()** operation

```
signal(S) {  
    S++;  
}
```

Binary semaphore:  
When s is 0 or 1, it is  
a mutex lock

# wail() and signal()

- Use board



# Semaphore Usage

- **Counting semaphore** – integer value can range over an unrestricted domain
- **Binary semaphore** – integer value can range only between 0 and 1
  - Same as a **mutex lock**
- Can solve various synchronization problems
- Consider  $P_1$  and  $P_2$  that require  $S_1$  to happen before  $S_2$   
Create a semaphore “**synch**” initialized to 0

P1 :

$S_1$  ;

**signal (synch) ;**

P2 :

**wait (synch) ;**

$S_2$  ;

- Can implement a counting semaphore  $S$  as a binary semaphore

# The counting semaphore

- **Controls access to a finite set of resources**
- Initialized to the number of resources
- Usage:
  - Wait (S): to use a resource
  - Signal (S): to release a resource
- When all resources are being used:  $S == 0$ 
  - Block until  $S > 0$  to use the resource

# Semaphore Implementation

- Must guarantee that no two processes can execute the `wait()` and `signal()` on the same semaphore at the same time
- Thus, the implementation becomes the critical section problem where the `wait` and `signal` code are placed in the critical section
  - Could now have **busy waiting** in critical section implementation
    - But implementation code is short
    - Little busy waiting if critical section rarely occupied
- Note that applications may spend lots of time in critical sections and therefore this is not a good solution
- Alternative: block and wakeup (next slide)

# Semaphore Implementation with no Busy waiting

- With each semaphore there is an associated waiting queue
- Each entry in a waiting queue has two data items:
  - value (of type integer)
  - pointer to next record in the list
- Two operations:
  - **block** – place the process invoking the operation on the appropriate waiting queue
  - **wakeup** – remove one of processes in the waiting queue and place it in the ready queue
- ```
typedef struct{  
    int value;  
    struct process *list;  
} semaphore;
```

## Implementation with no Busy waiting (Cont.)

```
wait(semaphore *S) {  
    S->value--;  
    if (S->value < 0) {  
        add this process to S->list;  
        block();  
    }  
}  
  
signal(semaphore *S) {  
    S->value++;  
    if (S->value <= 0) {  
        remove a process P from S->list;  
        wakeup(P);  
    }  
}
```

If value < 0  
abs(value) is the number  
of waiting processes

```
typedef struct{  
    int value;  
    struct process *list;  
} semaphore;
```

# Deadlock and Starvation

- **Deadlock** – two or more processes are waiting indefinitely for an event that can be caused by only one of the waiting processes
- Let  $s$  and  $q$  be two semaphores initialized to 1

$P_0$

```
wait(S) ;  
wait(Q) ;  
...  
signal(S) ;  
signal(Q) ;
```

$P_1$

```
wait(Q) ;  
wait(S) ;  
...  
signal(Q) ;  
signal(S) ;
```

- **$P_0$  executes wait(s),  $P_1$  executes wait(Q)**
  - $P_0$  must wait till  $P_1$  executes signal(Q)
  - $P_1$  must wait till  $P_0$  executes signal(S)      Deadlock!

# Priority Inversion

- **Priority Inversion** – Scheduling problem when lower-priority process holds a lock needed by higher-priority process
- Solved via **priority-inheritance protocol**
  - Process accessing resource needed by higher priority process Inherits higher priority till it finishes resource use
  - Once done, process reverts to lower priority

# Classical Problems of Synchronization

- Classical problems used to test newly-proposed synchronization schemes
  - Bounded-Buffer Problem
  - Readers and Writers Problem
  - Dining-Philosophers Problem



# Bounded-Buffer Problem

- $n$  buffers, each can hold one item
- Binary semaphore (mutex)
  - Provides mutual exclusion for accesses to buffer pool
  - Initialized to 1
- Counting semaphores
  - empty: Number of empty slots available
    - Initialized to  $n$
  - full: Number of filled slots available  $n$ 
    - Initialized to 0

# Bounded-Buffer : Note

- Producer and consumer must be ready before they attempt to enter critical section
- Producer readiness?
  - When a slot is available to add produced item
    - wait(empty): empty is initialized to n
- Consumer readiness?
  - When a producer has added new item to the buffer
    - wait(full) : full initialized to 0

# Bounded Buffer Problem (Cont.)

## The structure of the producer process

```
do {  
    ...  
    /* produce an item in next_produced */  
    ...  
    wait(empty);           wait till slot available  
    wait(mutex);          Allow producer OR consumer to (re)enter critical section  
    ...  
    /* add next produced to the buffer */  
    ...  
    signal(mutex);        Allow producer OR consumer to (re)enter critical section  
    signal(full);          signal consumer that a slot is available  
} while (true);
```

# Bounded Buffer Problem (Cont.)

## The structure of the consumer process

```
Do {  
    wait(full); wait till slot available for consumption  
    wait(mutex); Only producer OR consumer can be in critical section  
    ...  
    /* remove an item from buffer to next_consumed */  
    ...  
    signal(mutex); Allow producer OR consumer to (re)enter critical section  
    signal(empty); signal producer that a slot is available to add  
    ...  
    /* consume the item in next consumed */  
    ...  
} while (true);
```

# Notes

- Midterm Friday
- Background, Processes/Threads, Scheduling, Process Synchronization (exclude Classical Problems)
- Homework due Tuesday
  - No late period, solution shared Wednesday
- PA 3 available
  - Extension of PA2: interaction using pipes and shared memory
- Help session Thursday: Midterm, PA3

# Readers-Writers Problem

- A data set is shared among a number of concurrent processes
  - Readers – only read the data set; they do **not** perform any updates
  - Writers – can both read and write
- Problem – allow multiple readers to read at the same time
  - Only one single writer can access the shared data at the same time
- Several variations of how readers and writers are considered – all involve some form of priorities
- Shared Data
  - Data set
  - Semaphore `rw_mutex` initialized to 1 (mutual exclusion for writer)
  - Semaphore `mutex` initialized to 1 (mutual exclusion for `read_count`)
  - Integer `read_count` initialized to 0 (how many readers?)

# Readers-Writers Problem (Cont.)

- The structure of a writer process

```
do {  
    wait(rw_mutex);  
    ...  
    /* writing is performed */  
    ...  
    signal(rw_mutex);  
} while (true);
```

# Readers-Writers Problem (Cont.)

- The structure of a reader process


```
do {  
    wait(mutex);  
    read_count++;  
    if (read_count == 1)  
        wait(rw_mutex);  
    signal(mutex);  
  
    ...  
    /* reading is performed */  
    ...  
    wait(mutex);  
    read_count--;  
    if (read_count == 0)  
        signal(rw_mutex);  
    signal(mutex);  
} while (true);
```



# Readers-Writers Problem Variations


- **First** variation – no reader kept waiting unless writer has permission to use shared object
- **Second** variation – once writer is ready, it performs the write ASAP
- Both may have starvation leading to even more variations
- Problem is solved on some systems by kernel providing reader-writer locks

# Dining-Philosophers Problem



- Philosophers spend their lives alternating thinking and eating
- Don't interact with their neighbors, occasionally try to pick up 2 chopsticks (one at a time) to eat from bowl
  - Need both to eat, then release both when done
- In the case of 5 philosophers
  - Shared data
    - Bowl of rice (data set)
    - Semaphore **chopstick** [5] initialized to 1

# Dining-Philosophers Problem



# Dining-Philosophers Problem Algorithm: Simple solution?

- The structure of Philosopher *i*:

```
do {  
    wait (chopstick[i] );  
    wait (chopstick[ (i + 1) % 5] );  
  
    // eat  
  
    signal (chopstick[i] );  
    signal (chopstick[ (i + 1) % 5] );  
  
    // think  
  
} while (TRUE);
```

- What is the problem with this algorithm?
  - If all of them pick up the the left chopstick first -  
Deadlock

- Deadlock handling
  - Allow at most 4 philosophers to be sitting simultaneously at the table (with the same 5 forks).
  - Allow a philosopher to pick up the forks only if both are available (picking must be done in a critical section).
  - Use an asymmetric solution -- an odd-numbered philosopher picks up first the left chopstick and then the right chopstick. Even-numbered philosopher picks up first the right chopstick and then the left chopstick.

# Problems with Semaphores

- Incorrect use of semaphore operations:
  - `signal (mutex) .... wait (mutex)`: what happens?
  - `wait (mutex) ... wait (mutex) )`: what happens?
  - Omitting of `wait (mutex)` or `signal (mutex)` (or both): what happens?
- Deadlock and starvation are possible.