

# CS370 Operating Systems

Colorado State University

Yashwant K Malaiya

Fall 2016 Lecture 21



## Slides based on

- Text by Silberschatz, Galvin, Gagne
- Various sources

# FAQ

- Why not use a Boolean variable instead of Mutex?
- If there are more than one processes waiting, will the semaphore value be negative?
  - Negative: number of processes/threads waiting
  - 0: no waiting threads
  - Positive: no waiting threads, a wait operation would not put in queue the invoking thread. Often +1
- How to keep a philosopher from starving?
  - There exist solutions that will avoid a deadlock. However they may allow starvation, unless solution is further refined.
- Why not give each philosopher 2 chopsticks?
  - Nice and elegant solution. Widely used in Chinese restaurants. But takes all the fun away from the problem.

# FAQ

- Producer-consumer with bounded buffer
  - Should the production and consumption rates be a perfect match?
  - Can the producer add more than 1 item at a time?
- Monitors: what are they and how to implement them.
  - Details coming up.

# Classical Problems of Synchronization

- Classical problems used to test newly-proposed synchronization schemes
  - Bounded-Buffer Problem
  - Readers and Writers Problem
  - Dining-Philosophers Problem
- Monitors

# Notes

- PA 3 due Friday 11:59 PM
- Project topics and team approach: on canvas now

# Notes

## Research resources

- Books/Web articles
- Technical news
- [IEEE Explore](#)
- [ACM Digital Library](#)
- [ScienceDirect](#)

[Accessing library resources from Home](#)

[Google Scholar](#)

# Problems with Semaphores

- Incorrect use of semaphore operations:
  - Omitting of wait (mutex)
    - Violation of mutual exclusion
  - or signal (mutex)
    - Deadlock!

# Monitors

- Monitor: A high-level abstraction that provides a convenient and effective mechanism for process synchronization
- *Abstract data type*, internal variables only accessible by code within the procedure
- Only one process may be active within the monitor at a time
  - Automatically provide mutual exclusion
- Originally proposed for Concurrent Pascal 1975
- Directly supported by Java but not C




# Monitors

```
monitor monitor-name
{
    // shared variable declarations
    procedure P1 (...) { ... }

    procedure Pn (...) {.....}

    Initialization code (...) { ... }
}
}
```

# Schematic view of a Monitor



Only one process/thread in the Monitor

Provides an easy way to achieve mutual exclusion

But ... we also need a way for processes to **block** when they cannot proceed

# Condition Variables

The **condition** construct


Compare with  
semaphore

- **condition** **x**, **y**;
- Two operations are allowed on a condition variable:
  - **x.wait()** – a process that invokes the operation is suspended until **x.signal()**
  - **x.signal()** – resumes one of processes (if any) that invoked **x.wait()**
    - If no **x.wait()** on the variable, then it has no effect on the variable

## Difference between the signal() in semaphores and monitors

- Condition variables in Monitors: Not persistent
  - If a signal is performed and no waiting threads?
    - Signal is simply ignored
  - During subsequent wait operations
    - Thread blocks
- Semaphores
  - Signal increments semaphore value even if there are no waiting threads
    - Future wait operations would immediately succeed!

# Monitor with Condition Variables



# Condition Variables Choices

- If process P invokes `x.signal()`, and process Q is suspended in `x.wait()`, what should happen next?
  - Both Q and P cannot execute in parallel. If Q is resumed, then P must wait
- Options include
  - **Signal and wait** – P waits until Q either leaves the monitor or it waits for another condition
  - **Signal and continue** – Q waits until P either leaves the monitor or it waits for another condition
  - Both have pros and cons – language implementer can decide
  - Monitors implemented in [Concurrent Pascal \('75\)](#) compromise
    - P executing signal immediately leaves the monitor, Q is resumed
  - Implemented in other languages including C#, Java

## Difference between the signal() in semaphores and monitors

- Monitors {condition variables}: Not persistent
  - If a signal is performed and no waiting threads?
    - Signal is simply ignored
  - During subsequent wait operations
    - Thread blocks
- Semaphores
  - Signal increments semaphore value even if there are no waiting threads
  - Future wait operations would immediately succeed!

- Research in distributed systems and predictive analytics
  - Big data, virtualized cloud
- CS455: Introduction to Distributed Systems
  - Spring 2017



# Monitor Solution to Dining Philosophers: Deadlock-free

```
enum {THINKING, HUNGRY, EATING} state[5];
```

- `state[i] = EATING` only if
  - `state[(i+4)%5] != EATING && state[(i+1)%5] != EATING !`
- `condition self[5]`
  - Delay self when **HUNGRY but unable** to get chopsticks

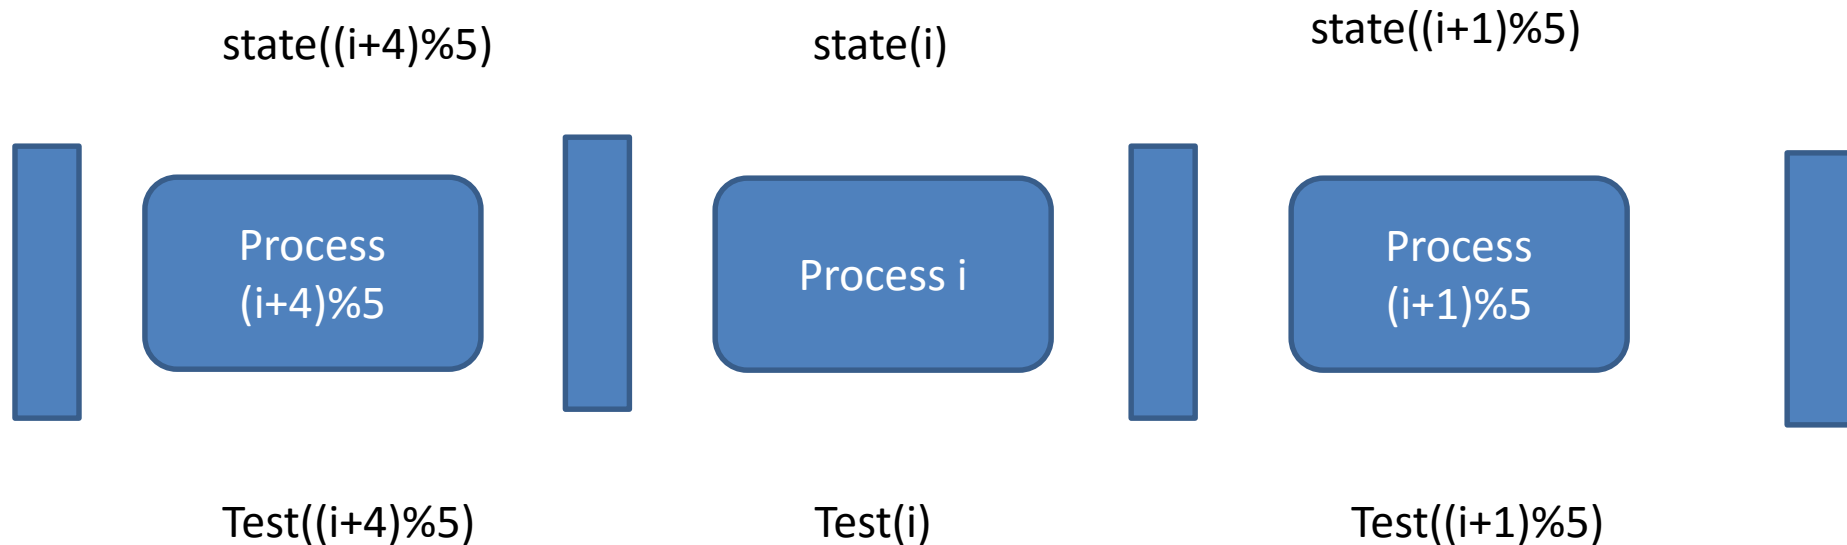
## Sequence of actions

- Before eating, must invoke `pickup()`
  - May result in suspension of philosopher process
  - After completion of operation, philosopher may eat

```
think  
DiningPhilosophers.pickup(i);  
eat  
DiningPhilosophers.putdown(i);  
think
```

# Monitor Solution to Dining Philosophers: Deadlock-free

```
enum { THINKING, HUNGRY, EATING } state[5];
```



# The pickup() and putdown() operations

```
monitor DiningPhilosophers
```

```
{
    enum { THINKING; HUNGRY, EATING) state [5] ;
    condition self [5];

    void pickup (int i) {
        state[i] = HUNGRY;
        test(i);    //on next slide
        if (state[i] != EATING) self[i].wait;
    }

    void putdown (int i) {
        state[i] = THINKING;
        // test left and right neighbors
        test((i + 4) % 5);
        test((i + 1) % 5);
    }
}
```

Suspend self if  
unable to acquire  
chopstick

Check to see if person  
on left or right can use  
the chopstick



# test() to see if philosopher I can eat



Eat only if HUNGRY  
and Person on Left  
AND Right  
are not eating

```
void test (int i) {  
    if ((state[(i + 4) % 5] != EATING) &&  
        (state[i] == HUNGRY) &&  
        (state[(i + 1) % 5] != EATING) ) {  
        state[i] = EATING ;  
        self[i].signal () ;  
    }  
}
```

Signal a process that  
was suspended while  
trying to eat

```
initialization_code() {  
    for (int i = 0; i < 5; i++)  
        state[i] = THINKING;  
}
```

# Possibility of starvation

- Philosopher  $i$  can starve if eating periods of philosophers on left and right overlap
- Possible solution
  - Introduce new state: STARVING
  - Chopsticks can be picked up if no neighbor is starving
    - Effectively wait for neighbor's neighbor to stop eating
    - REDUCES concurrency!

# Monitor Implementation Using Semaphores

# Monitor Implementation Mutual Exclusion

For each monitor

- Semaphore mutex initialized to 1
- Process must execute
  - wait(mutex) : Before entering the monitor
  - signal(mutex): Before leaving the monitor

# Monitor Implementation Using Semaphores

- Variables

```
semaphore mutex; // (initially = 1) allows only one process to be active
semaphore next;  // (initially = 0) causes signaler to sleep
int next_count = 0; num of sleepers since they signalled
```

- Each procedure  $F$  will be replaced the compiler by

```
wait(mutex);
...
body of F;
...
if (next_count > 0)
    signal(next)
else
    signal(mutex);
```

Both mutex and next have  
an associated queue

- Mutual exclusion within a monitor is ensured