

CS370 Operating Systems

Colorado State University

Yashwant K Malaiya

Fall 2016 Lecture 22



Slides based on

- Text by Silberschatz, Galvin, Gagne
- Various sources

FAQ

- Classes that follow CS370
 - CS455 Distributed Systems Spring
 - CS457 Networks Fall
 - CS470 Computer Architecture Spring
 - CS475 Parallel Programming Fall
 - CS435: Introduction to Big Data Spring

Midterm

- Scores available Friday
 - Raw -> Adjusted
 - Linear: $y = mx + c$ $m = 0.703$, $c = 35.239$
 - Adjusted scores:
 - Max 102
 - Average 80
- Tests handed back on Monday


Condition Variables

The **condition** construct

Compare with
semaphore

- **condition** **x**, **y**;
- Two operations are allowed on a condition variable:
 - **x.wait()** – a process that invokes the operation is suspended until **x.signal()**
 - **x.signal()** – resumes one of processes (if any) that invoked **x.wait()**
 - If no **x.wait()** on the variable, then it has no effect on the variable

Monitor with Condition Variables



Condition Variables Choices

- If process P invokes `x.signal()`, and process Q is suspended in `x.wait()`, what should happen next?
 - Both Q and P cannot execute in parallel. If Q is resumed, then P must wait
- Options include
 - **Signal and wait** – P waits until Q either leaves the monitor or it waits for another condition
 - **Signal and continue** – Q waits until P either leaves the monitor or it waits for another condition
 - Both have pros and cons – language implementer can decide
 - Monitors implemented in [Concurrent Pascal \('75\)](#) compromise
 - P executing signal immediately leaves the monitor, Q is resumed
 - Implemented in other languages including C#, Java

Monitor Solution to Dining Philosophers: Deadlock-free

```
enum {THINKING, HUNGRY, EATING} state[5];
```

- `state[i] = EATING` only if
 - `state[(i+4)%5] != EATING && state[(i+1)%5] != EATING !`
- `condition self[5]`
 - Delay self when **HUNGRY but unable** to get chopsticks

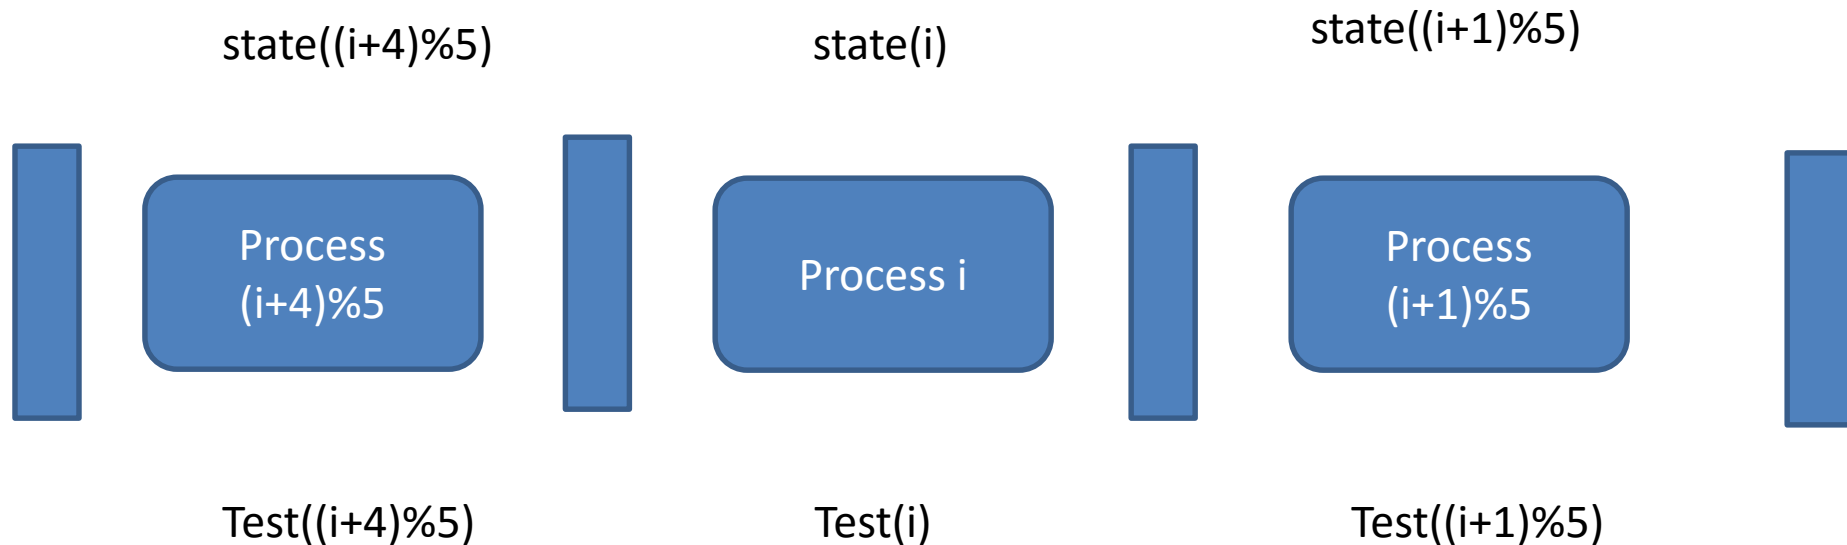
Sequence of actions

- Before eating, must invoke `pickup()`
 - May result in suspension of philosopher process
 - After completion of operation, philosopher may eat

```
think  
DiningPhilosophers.pickup(i);  
eat  
DiningPhilosophers.putdown(i);  
think
```

Monitor Solution to Dining Philosophers: Deadlock-free

```
enum { THINKING, HUNGRY, EATING } state[5];
```



The pickup() and putdown() operations

```
monitor DiningPhilosophers
```

```
{  
    enum { THINKING; HUNGRY, EATING) state [5] ;  
    condition self [5];  
  
    void pickup (int i) {  
        state[i] = HUNGRY;  
        test(i);    //on next slide  
        if (state[i] != EATING) self[i].wait;  
    }  
}
```

Suspend self if
unable to acquire
chopstick

```
    void putdown (int i) {  
        state[i] = THINKING;  
        // test left and right neighbors  
        test((i + 4) % 5);  
        test((i + 1) % 5);  
    }  
}
```

Check to see if person
on left or right can use
the chopstick



test() to see if philosopher I can eat



Eat only if HUNGRY
and Person on Left
AND Right
are not eating

```
void test (int i) {  
    if ((state[(i + 4) % 5] != EATING) &&  
        (state[i] == HUNGRY) &&  
        (state[(i + 1) % 5] != EATING) ) {  
        state[i] = EATING ;  
        self[i].signal () ;  
    }  
}
```

Signal a process that
was suspended while
trying to eat

```
initialization_code() {  
    for (int i = 0; i < 5; i++)  
        state[i] = THINKING;  
}
```

Possibility of starvation

- Philosopher i can starve if eating periods of philosophers on left and right overlap
- Possible solution
 - Introduce new state: STARVING
 - Chopsticks can be picked up if no neighbor is starving
 - Effectively wait for neighbor's neighbor to stop eating
 - REDUCES concurrency!

Monitor Implementation Using Semaphores

Monitor Implementation Mutual Exclusion

For each monitor

- Semaphore mutex initialized to 1
- Process must execute
 - wait(mutex) : Before entering the monitor
 - signal(mutex): Before leaving the monitor

Monitor Implementation Using Semaphores

- Variables

```
semaphore mutex;    // (initially = 1) allows only one process to be active
semaphore next;     // (initially = 0) causes signaler to sleep
int next_count = 0;  // num of sleepers since they signalled
```

- Each procedure **F** will be replaced the compiler by

```
wait(mutex) ;
...
body of F;
...
if (next_count > 0)
    signal(next)
else
    signal(mutex) ;
```

Both mutex and next have
an associated queue

- Mutual exclusion within a monitor is ensured

Monitor Implementation – Condition Variables

- For each condition variable x , we have:

```
semaphore x_sem; // (initially = 0) causes caller of  
                  wait to sleep  
int x_count = 0; // number of sleepers on condition
```

- The operations **x .wait** and **x .signal** can be implemented as:

The operation x .wait can be implemented as:

```
x_count++;  
if (next_count > 0)  
    signal(next);  
else  
    signal(mutex);  
wait(x_sem);  
x_count--;
```

The operation x .signal can be implemented as:

```
if (x_count > 0) {  
    next_count++;  
    signal(x_sem);  
    wait(next);  
    next_count--;  
}
```

Resuming Processes within a Monitor

- If several processes queued on condition x , and $x.\text{signal}()$ is executed, which should be resumed?
- FCFS frequently not adequate
- **conditional-wait** construct of the form $x.\text{wait}(c)$
 - Where c is **priority number**
 - Process with lowest number (highest priority) is scheduled next

Single Resource allocation

- Allocate a single resource among competing processes using priority numbers that specify the maximum time a process plans to use the resource

```
R.acquire(t) ;  
    ...  
    access the resource ;  
    ...  
R.release ;
```

- Where R is an instance of type **ResourceAllocator**

A Monitor to Allocate Single Resource

```
monitor ResourceAllocator
```

```
{
```

```
    boolean busy;
```

```
    condition x;
```

```
    void acquire(int time) {
```

```
        if (busy)
```

```
            x.wait(time);
```

```
        busy = TRUE;
```

```
    }
```

```
    void release() {
```

```
        busy = FALSE;
```

```
        x.signal();
```

```
    }
```

```
    initialization code() {
```

```
        busy = FALSE;
```

```
    }
```

```
}
```

Sleep, Time used
to prioritize
waiting
processes

Wakes up
one of the
processes

Java Synchronization

- For simple synchronization Java provides the synchronized keyword
 - synchronizing methods
`public synchronized void increment() { c++; }`
 - synchronizing blocks
`synchronized(this) {
 lastName = name;
 nameCount++;
}`
- wait() and notify() allows a thread to wait for an event. A call to notify. all() allows all threads that are on wait() with the same lock to be released
- For more sophisticated locking mechanisms, starting from Java 5, the package java.concurrent.locks provides additional locking

Synchronization Examples

- Solaris
- Windows
- Linux
- Pthreads

Solaris Synchronization

- Implements a variety of locks to support multitasking, multithreading (including real-time threads), and multiprocessing
- Uses **adaptive mutexes** for efficiency when protecting data from short code segments
 - Starts as a standard semaphore spin-lock
 - If lock held, and by a thread running on another CPU, spins
 - If lock held by non-run-state thread, block and sleep waiting for signal of lock being released
- Uses **condition variables**
- Uses **readers-writers** locks when longer sections of code need access to data
- Uses **turnstiles** to order the list of threads waiting to acquire either an adaptive mutex or reader-writer lock
 - Turnstiles are per-lock-holding-thread, not per-object
- Priority-inheritance per-turnstile gives the running thread the highest of the priorities of the threads in its turnstile

Windows Synchronization

- Uses interrupt masks to protect access to global resources on uniprocessor systems
- Uses **spinlocks** on multiprocessor systems
 - Spinlocking-thread will never be preempted
- Also provides **dispatcher objects** user-land which may act mutexes, semaphores, events, and timers
 - **Events**
 - An event acts much like a condition variable
 - Timers notify one or more thread when time expired
 - Dispatcher objects either **signaled-state** (object available) or **non-signaled state** (thread will block)

Linux Synchronization

- Linux:
 - Prior to kernel Version 2.6, disables interrupts to implement short critical sections
 - Version 2.6 and later, fully preemptive
- Linux provides:
 - Semaphores
 - atomic integers
 - spinlocks
 - reader-writer versions of both
- On single-cpu system, spinlocks replaced by enabling and disabling kernel preemption

Pthreads Synchronization

- Pthreads API is OS-independent
- It provides:
 - mutex locks
 - condition variable
- Non-portable extensions include:
 - read-write locks
 - spinlocks

Alternative Approaches

- Transactional Memory
- OpenMP
- Functional Programming Languages

Transactional Memory

- A **memory transaction** is a sequence of read-write operations to memory that are performed atomically.

```
void update()  
{  
    /* read/write memory */  
}
```

OpenMP

- OpenMP is a set of compiler directives and API that support parallel programming.

```
void update(int value)
{
    #pragma omp critical
    {
        count += value
    }
}
```

The code contained within the `#pragma omp critical` directive is treated as a critical section and performed atomically.