

CS370 Operating Systems

Colorado State University

Yashwant K Malaiya

Fall 2016 Lecture 24



Deadlocks

Slides based on

- Text by Silberschatz, Galvin, Gagne
- Various sources

FAQ

- Resource allocation graph: what are resources?
Memory space, IO devices, files, etc.
- Deadlock definition: A set of processes is deadlocked if each process in the set is waiting for an event that only another process in the set can cause
- Deadlock Characterization: 4 *necessary* conditions
- Edges: request $P \rightarrow R$ assignment $R \rightarrow P$
- If a cycle is present:
 - If each resource type has exactly one instance: Deadlock **has** occurred
 - If each resource type has multiple instances: A deadlock **may** have occurred

Notes

- Project topic, team, coordinator:
 - [submission](#) due today
 - Soon: select a group name and join the name
- PA4 Due 10/27. [Early reward due 10/20](#)
 - Note unlimited number of processes (use ArrayList).
Also consider Scanner.
- Java Monitor example (Dining Philosophers)
[available](#) as a self-exercise. See Piazza post.

Chapter 7: Deadlocks


- System Model
- Deadlock Characterization
- Methods for Handling Deadlocks
 - Deadlock Prevention
 - Deadlock Avoidance resource-allocation
 - Deadlock Detection
 - Recovery from Deadlock

Deadlock Characterization

Deadlock **can** arise if four conditions hold simultaneously.

- **Mutual exclusion**: only one process at a time can use a resource
- **Hold and wait**: a process holding at least one resource is waiting to acquire additional resources held by other processes
- **No preemption**: a resource can be released only voluntarily by the process holding it, after that process has completed its task
- **Circular wait**: there exists a set $\{P_0, P_1, \dots, P_n\}$ of waiting processes such that P_0 is waiting for a resource that is held by P_1 , P_1 is waiting for a resource that is held by P_2 , ..., P_{n-1} is waiting for a resource that is held by P_n , and P_n is waiting for a resource that is held by P_0 .

Graph With A Cycle But No Deadlock



There is no deadlock. P_4 may release its instance of resource type R_2 . That resource can then be allocated to P_3 , breaking the cycle

If a resource-allocation graph does not have a cycle, then the system is **not** in a deadlocked state.

If there is a cycle, then the system may or may not be in a deadlocked state.

Basic Facts

- If graph contains no cycles \Rightarrow no deadlock
- If graph contains a cycle \Rightarrow
 - if only one instance per resource type, then deadlock
 - if several instances per resource type, possibility of deadlock

Methods for Handling Deadlocks

- Ensure that the system will *never* enter a deadlock state:
 - Deadlock prevention
 - ensuring that at least one of the 4 conditions cannot hold
 - Deadlock avoidance
 - Dynamically examines the resource-allocation state to ensure that there can never be a circular-wait condition
- Allow the system to enter a deadlock state
 - Detect and then recover
- Ignore the problem and pretend that deadlocks never occur in the system; used by most operating systems, including UNIX

Methods for Handling Deadlocks

- **Deterministic**: Ensure that the system will *never* enter a deadlock state at any cost
- Handle if it happens: Allow the system to enter a deadlock state and then recover
- Ostrich algorithm: Stick your head in the sand; pretend there is no problem at all .
 - May be acceptable if it happens only rarely
(**Probabilistic view**)

Ostrich algorithm

Advantages:

- Cheaper, rarely needed anyway
- Prevention, avoidance, detection and recovery
 - Need to run constantly

Disadvantages:

- Resources held by processes that cannot run
- More and more processes enter deadlocked state
 - When they request more resources
- Deterioration in system performance
 - Requires restart

Deadlock Prevention: Limit Mutual Exclusion

For a deadlock to occur, each of the four necessary conditions must hold. By ensuring that at least one of these conditions cannot hold, we can **prevent** the occurrence of a deadlock.

Restrain the ways request can be made:

- **Limit Mutual Exclusion** –
 - not required for sharable resources (e.g., read-only files)
 - Mutual Exclusion must hold for non-sharable resources



Deadlock Prevention: Limit Hold and Wait

- **Limit Hold and Wait** – must guarantee that whenever a process requests a resource, it does not hold any other resources
 1. Require process to request and be allocated all its resources before it begins execution



Deadlock Prevention: Limit Hold and Wait

- **Limit Hold and Wait** – must guarantee that whenever a process requests a resource, it does not hold any other resources
 2. Allow a process to request resources when it is holding none.
Ex: Copy data from DVD, sort file, and print
 - First request DVD and disk: file
 - Then request file and printer
- Disadvantage: starvation possible



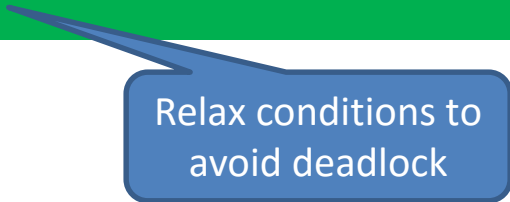
Deadlock Prevention: Limit No Preemption

- **Limit No Preemption** –
 - If a process that is holding some resources, requests another resource that cannot be immediately allocated to it, then all resources currently being held are released
 - *Preempted resources* are added to the list of resources for which the process is waiting
 - Process will be restarted only when it can regain its old resources, as well as the new ones that it is requesting

Deadlock Prevention: Limit Circular Wait

- **Limit Circular Wait** – impose a total ordering of all resource types, and require that each process requests resources in an increasing order of enumeration
- Assign each resource a unique number
 - Disk drive: 1
 - Printer: 2 ...
 - Request resources in increasing order

Dining philosophers problem: Necessary conditions for deadlock



Relax conditions to avoid deadlock

- Mutual exclusion
 - 2 philosophers *cannot share* the same chopstick
- Hold-and-wait
 - A philosopher *picks up one* chopstick at a time
 - Will not let go of the first while it *waits for the second* one
- No preemption
 - A philosopher *does not snatch chopsticks* held by some other philosopher
- Circular wait
 - Could happen if each philosopher *picks chopstick with the same hand* first

Deadlock Example

```
/* thread one runs in this function */
void *do_work_one(void *param)
{
    pthread_mutex_lock(&first_mutex);
    pthread_mutex_lock(&second_mutex);
    /** * Do some work */
    pthread_mutex_unlock(&second_mutex);
    pthread_mutex_unlock(&first_mutex);
    pthread_exit(0);
}

/* thread two runs in this function */
void *do_work_two(void *param)
{
    pthread_mutex_lock(&second_mutex);
    pthread_mutex_lock(&first_mutex);
    /** * Do some work */
    pthread_mutex_unlock(&first_mutex);
    pthread_mutex_unlock(&second_mutex);
    pthread_exit(0);
}
```

Assume that thread one is the first to acquire the locks and does so in the order (1) first mutex, (2) second mutex.

Solution: Lock-order verifier, **Witness** records the relationship that **first mutex must be acquired before second mutex**. If thread two later acquires the locks out of order, witness generates a warning message on the system console.

Deadlock Example with Lock Ordering

```
void transaction(Account from, Account to, double amount)
{
    mutex lock1, lock2;
    lock1 = get_lock(from);
    lock2 = get_lock(to);
    acquire(lock1);
    acquire(lock2);
    withdraw(from, amount);
    deposit(to, amount);
    release(lock2);
    release(lock1);
}
```

Lock ordering:
First from lock, then to lock

Ex: Transactions 1 and 2 execute concurrently.

Transaction 1 transfers \$25 from account A to account B, and

Transaction 2 transfers \$50 from account B to account A.

Deadlock is possible, even with lock ordering.

Deadlock Avoidance

Deadlock Avoidance

- Require additional information about how resources are to be requested
- Knowledge about sequence of requests and releases for processes
 - Allows us to decide if resource allocation could cause a **future deadlock**
 - Process P: Tape drive, then printer
 - Process Q: Printer, then tape drive

Deadlock Avoidance: Handling resource requests

- For each resource request:
 - Decide whether or not process should wait
 - To avoid possible future deadlock
- Predicated on:
 1. Currently available resources
 2. Currently allocated resources
 3. *Future requests and releases of each process*

Avoidance: amount and type of information needed

- **Resource allocation state**
 - Number of available and allocated resources
 - Maximum demands of processes
- *Dynamically* examine resource allocation state
 - Ensure circular-wait cannot exist
- Simplest model:
 - Declare maximum number of resources for each type
 - Use information to avoid deadlock

Safe Sequence

System must decide if immediate allocation leaves the system in a safe state


System is in **safe state** if there exists a sequence $\langle P_1, P_2, \dots, P_n \rangle$ of ALL the processes such that

- for each P_i , the resources that P_i can still request can be satisfied by
 - currently available resources +
 - resources held by all the P_j , with $j < i$
 - That is
 - If P_i resource needs are not immediately available, then P_i can wait until all P_j have finished and released resources
 - When P_i terminates, P_{i+1} can obtain its needed resources, and so on
- If no such sequence exists: system state is **unsafe**

Deadlock avoidance: Safe states

- If the system can:
 - Allocate resources to each process in some order
 - Up to the maximum for the process
 - Still avoid deadlock
 - Then it is in a **safe state**
- A system is safe ONLY IF there is a safe sequence
- A safe state is not a deadlocked state
 - Deadlocked state is an unsafe state
 - Not all unsafe states are deadlock

Safe, Unsafe, Deadlock State



A unsafe state may lead to deadlock