

# CS370 Operating Systems

Colorado State University

Yashwant K Malaiya

Fall 2016 Lecture 25



## Deadlock

Slides based on

- Text by Silberschatz, Galvin, Gagne
- Various sources

# FAQ

- Does a cycle in a resource allocation graph signify “circular wait”? Only if there is only one instance of a resource
- Safe state idea: What about resources held by processes that have already run? they have been released and are thus available.
- Does the MAC spinning wheel indicate a deadlock?
- Are application hangs caused by deadlocks? OS timer perhaps 5 sec

# Chapter 7: Deadlocks

- System Model
- Deadlock Characterization
- Methods for Handling Deadlocks
  - Deadlock Prevention
  - Deadlock Avoidance
  - Deadlock Detection
  - Recovery from Deadlock

# Deadlock Avoidance

Requires that the system has some additional *a priori* information available

- Simplest and most useful model requires that each process declare the *maximum number* of resources of each type that it may need
- The deadlock-avoidance algorithm dynamically examines the resource-allocation state to ensure that there can never be a circular-wait condition
- Resource-allocation *state* is defined by the number of available and allocated resources, and the maximum demands of the processes


# Safe State

- When a process requests an available resource, system must decide if immediate allocation leaves the system in a safe state
- System is in **safe state** if there exists a sequence  $\langle P_1, P_2, \dots, P_n \rangle$  of ALL the processes in the systems such that for each  $P_i$ , the resources that  $P_i$  can still request can be satisfied by currently available resources + resources held by all the  $P_j$ , with  $j < i$
- That is:
  - If  $P_i$  resource needs are not immediately available, then  $P_i$  can wait until all  $P_j$  have finished
  - When  $P_j$  is finished,  $P_i$  can obtain needed resources, execute, return allocated resources, and terminate
  - When  $P_i$  terminates,  $P_{i+1}$  can obtain its needed resources, and so on

# Basic Facts

- If a system is in safe state  $\Rightarrow$  no deadlocks
- If a system is in unsafe state  $\Rightarrow$  possibility of deadlock
- Avoidance  $\Rightarrow$  ensure that a system will never enter an unsafe state.

# Safe, Unsafe, Deadlock State



## Example A: 12 Units available in the system

	Max need	Current holding
P0	10	5
P1	4	2
P2	9	2

**At T0:**

9 units allocated

3 units available

*A unit could be a drive,  
a block of memory etc.*

- At time **T0** is the system is in a safe state?
  - Try sequence  $\langle P1, P0, P2 \rangle$
  - P1 can be given 2 units
  - When P1 releases its resources; there are 5 units
  - P0 uses 5 and subsequently releases them (# 10 now)
  - P2 can then proceed.
- Thus  $\langle P1, P0, P2 \rangle$  is a safe sequence, and at T0 system was in a safe state



## Example B: 12 Units available in the system

	Max need	Current holding
P0	10	5
P1	4	2
P2	9	2+1

**Before T1:**

3 units available

**At T1:**

2 units available

- At time **T1**, P2 is allocated 1 more units. Is that a good decision?
  - Now only P1 can proceed.
  - When P1 releases its resources; there are 4 units
  - P0 needs 5 and P2 needs 6. Deadlock.
    - **Mistake** in granting P2 additional units
- The state at **T1** is not a safe state.

# Avoidance Algorithms



- Single instance of a resource type
  - Use a resource-allocation graph scheme
- Multiple instances of a resource type
  - Use the banker's algorithm



# Resource-Allocation Graph Scheme

- **Claim edge**  $P_i \rightarrow R_j$  indicated that process  $P_i$  **may** request resource  $R_j$ ; represented by a dashed line
- Claim edge converts to **request edge** when a process **requests** a resource
- Request edge converted to an **assignment edge** when the resource is **allocated** to the process
- When a resource is released by a process, assignment edge reconverts to a claim edge
- Resources must be claimed *a priori* in the system

# Resource-Allocation Graph



Suppose  $P_2$  requests  $R_2$ . Although  $R_2$  is currently free, we cannot allocate it to  $P_2$ , since this action will create a cycle getting system in an unsafe state. If  $P_1$  requests  $R_2$ , and  $P_2$  requests  $R_1$ , then a deadlock will occur.

# Resource-Allocation Graph Algorithm

- Suppose that process  $P_i$  requests a resource  $R_j$
- The request can be granted only if converting the request edge to an assignment edge does not result in the formation of a cycle in the resource allocation graph

# Banker's Algorithm: examining a request

- Multiple instances of resources.
- Each process must a priori claim maximum use
- When a process requests a resource
  - it may have to wait ([resource request algorithm](#))
  - Request not granted if the resulting system state is unsafe ([safety algorithm](#))
- When a process gets all its resources it must return them in a finite amount of time
- [Modeled after a banker in a small town making loans](#)

# Data Structures for the Banker's Algorithm

Let  $n$  = number of processes, and  $m$  = number of resources types.

- **Available:** Vector of length  $m$ . If available  $[j] = k$ , there are  $k$  instances of resource type  $R_j$  available

## Processes vs resources:

- **Max:**  $n \times m$  matrix. If  $Max[i,j] = k$ , then process  $P_i$  **may request at most**  $k$  instances of resource type  $R_j$
- **Allocation:**  $n \times m$  matrix. If  $Allocation[i,j] = k$  then  $P_i$  is **currently allocated**  $k$  instances of  $R_j$
- **Need:**  $n \times m$  matrix. If  $Need[i,j] = k$ , then  $P_i$  **may need**  $k$  more instances of  $R_j$  to complete its task

$$Need[i,j] = Max[i,j] - Allocation[i,j]$$

# Safety Algorithm: Is System in safe state?

1. Let **Work** and **Finish** be vectors of length  $m$  and  $n$ , respectively. Initialize:

**Work** = **Available**

**Finish** [ $i$ ] = **false** for  $i = 0, 1, \dots, n-1$

2. Find a process  $i$  such that both:

(a) **Finish** [ $i$ ] = **false**

(b) **Need** <sub>$i$</sub>  ≤ **Work**

If no such  $i$  exists, go to step 4

3. **Work** = **Work** + **Allocation** <sub>$i$</sub>   
**Finish** [ $i$ ] = **true**  
go to step 2

4. If **Finish** [ $i$ ] == **true** for all  $i$ , then the system is in a safe state

**n** = number of processes,  
**m** = number of resources types  
**Need** <sub>$i$</sub> : additional res needed  
**Work**: res currently free  
**Finish** <sub>$i$</sub> : processes finished  
**Allocation** <sub>$i$</sub> : allocated to  $i$



## Resource-Request Algorithm for Process $P_i$

**Notation:**  $\mathbf{Request}_i$  = request vector for process  $P_i$ .

If  $\mathbf{Request}_i[j] = k$  then process  $P_i$  wants  $k$  instances of resource type  $R_j$

**Algorithm:** *Should the allocation request be granted?*

1. If  $\mathbf{Request}_i \leq \mathbf{Need}_i$ , go to step 2. Otherwise, raise **error condition**, since process has exceeded its maximum claim
2. If  $\mathbf{Request}_i \leq \mathbf{Available}$ , go to step 3. Otherwise  $P_i$  must wait, since resources are **not available**
3. **Is allocation safe?:** **Pretend** to allocate requested resources to  $P_i$  by modifying the state as follows:
  - $\mathbf{Available} = \mathbf{Available} - \mathbf{Request}_i$
  - $\mathbf{Allocation}_i = \mathbf{Allocation}_i + \mathbf{Request}_i$
  - $\mathbf{Need}_i = \mathbf{Need}_i - \mathbf{Request}_i$
  - If safe  $\Rightarrow$  the resources are allocated to  $P_i$
  - If unsafe  $\Rightarrow P_i$  must wait, and the old resource-allocation state is preserved.

# Example of Banker's Algorithm

- 5 processes  $P_0$  through  $P_4$ ;
- 3 resource types:  $A$  (10 instances),  $B$  (5 instances), and  $C$  (7 instances)
- Snapshot at time  $T_0$ :

	<u>Allocation</u>	<u>Max</u>	<u>Available</u>
	$A \ B \ C$	$A \ B \ C$	$A \ B \ C$
$P_0$	0 1 0	7 5 3	3 3 2
$P_1$	2 0 0	3 2 2	
$P_2$	3 0 2	9 0 2	
$P_3$	2 1 1	2 2 2	
$P_4$	0 0 2	4 3 3	

- Is it a safe state?

# Example (Cont.)

- The matrix ***Need*** is ***Max – Allocation***

	<u><i>Need</i></u>		
	<i>A</i>	<i>B</i>	<i>C</i>
$P_0$	7	4	3
$P_1$	1	2	2
$P_2$	6	0	0
$P_3$	0	1	1
$P_4$	4	3	1

Available		
<i>A</i>	<i>B</i>	<i>C</i>
3	3	2

- The system is in a safe state since the sequence  $\langle P_1, P_3, P_4, P_2, P_0 \rangle$  satisfies safety criteria (see next)

# Example Cont.

The system is in a safe state since the sequence  $\langle P_1, P_3, P_4, P_2, P_0 \rangle$  satisfies safety criteria, since:

	<u>Allocation</u>	<u>Need</u>	<u>Available</u>
	A B C	A B C	A B C
$P_0$	0 1 0	7 4 3	3 3 2
$P_1$	2 0 0	1 2 2	
$P_2$	3 0 2	6 0 0	
$P_3$	2 1 1	0 1 1	
$P_4$	0 0 2	4 3 1	

P1 can run since  
need  $\leq$  available

P1 run to completion. Available becomes  $[3\ 3\ 2] + [2\ 0\ 0] = [5\ 3\ 2]$

P3 run to completion. Available becomes  $[5\ 3\ 2] + [2\ 1\ 1] = [7\ 4\ 3]$

P4 run to completion. Available becomes  $[7\ 4\ 3] + [0\ 0\ 2] = [7\ 4\ 5]$

P2 run to completion. Available becomes  $[7\ 4\ 5] + [3\ 0\ 2] = [10\ 4\ 7]$

P0 run to completion. Available becomes  $[10\ 4\ 7] + [0\ 1\ 0] = [10\ 5\ 7]$

**Hence state above is safe**

## Ex: Assume now $P_1$ Requests (1,0,2)

- Check that Request  $\leq$  Available
  - $(1,0,2) \leq (3,3,2) \Rightarrow$  true. Check for safety after pretend allocation.  $P_1$  allocation would be  $(2\ 0\ 0) + (1\ 0\ 2)$

	<u>Allocation</u>	<u>Need</u>	<u>Available</u>
	A B C	A B C	A B C
$P_0$	0 1 0	7 4 3	2 3 0
$P_1$	3 0 2	0 2 0	
$P_2$	3 0 2	6 0 0	
$P_3$	2 1 1	0 1 1	
$P_4$	0 0 2	4 3 1	

- Executing safety algorithm shows that sequence  $\langle P_1, P_3, P_4, P_0, P_2 \rangle$  satisfies safety requirement. Yes, safe state.

# Additional Requests ..

- Given State is (previous slide)

	<u>Allocation</u>	<u>Need</u>	<u>Available</u>
	A B C	A B C	A B C
$P_0$	0 1 0	7 4 3	2 3 0
$P_1$	3 0 2	0 2 0	
$P_2$	3 0 2	6 0 0	
$P_3$	2 1 1	0 1 1	
$P_4$	0 0 2	4 3 1	

- P4 request for (3,3,0): cannot be granted - resources are not available.
- P0 request for (0,2,0): cannot be granted since the resulting state is unsafe.

# Deadlock Detection


- Allow system to enter deadlock state
- Detection algorithm
  - Single instance of each resource:
    - wait-for graph
  - Multiple instances:
    - detection algorithm (based on Banker's algorithm)
- Recovery scheme

# Single Instance of Each Resource Type


- Maintain **wait-for** graph (based on resource allocation graph)
  - Nodes are processes
  - $P_i \rightarrow P_j$  if  $P_i$  is waiting for  $P_j$
  - *Deadlock if cycles*
- Periodically invoke an algorithm that searches for a cycle in the graph. If there is a cycle, there exists a deadlock
- An algorithm to detect a cycle in a graph requires an order of  $n^2$  operations, where  $n$  is the number of vertices in the graph



# Resource-Allocation Graph and Wait-for Graph



Resource-Allocation Graph



Corresponding wait-for graph

3 cycles. Deadlock.

# Several Instances of a Resource Type

**Banker's algorithm:** Can requests by all process be satisfied?

- **Available:** A vector of length  $m$  indicates the number of available (currently free) resources of each type
- **Allocation:** An  $n \times m$  matrix defines the number of resources of each type currently allocated to each process
- **Request:** An  $n \times m$  matrix indicates the current request of each process. If **Request**  $[i][j] = k$ , then process  $P_i$  is requesting  $k$  more instances of resource type  $R_j$ .