

# CS370 Operating Systems

Colorado State University

Yashwant K Malaiya

Fall 2016 Lecture 29



## Main Memory


Slides based on

- Text by Silberschatz, Galvin, Gagne
- Various sources

# FAQ

## Terminology:

- Contiguous allocation: Allocation of a contiguous memory region to a process (from address  $x$  to  $x+y$ )
- Internal vs external fragmentation
  - Internal fragmentation: memory wasted within an allocated memory region
  - External fragmentation: memory wasted due to small chunks of free memory interspersed among allocated regions



# FAQ

Paging:

Can we use really small pages/frames to minimize internal fragmentation within frames?

- Answer coming up

Paging and Addresses

- Part of the address identifies a page
- The other part identified location within the frame
- More coming up

# Pages

- Pages and frames
- Page tables
- TLB: page table caching
- Memory protection and sharing
- Multilevel page tables

# Paging


- Divide physical memory into fixed-sized blocks called **frames**
  - Size is power of 2, between 512 bytes and 16 Mbytes
- Divide logical memory into blocks of same size called **pages**
- To run a program of size  **$N$**  pages, need to find  **$N$**  free frames and load program
- Still have Internal fragmentation
- Physical address space of a process can be noncontiguous; process is allocated physical memory whenever the latter is available
  - Avoids external fragmentation
  - Avoids problem of varying sized memory chunks

# Paging

- physical memory - **frames**
- logical memory - **pages**
- Keep track of all free frames
- Set up a **page table** to translate logical to physical addresses

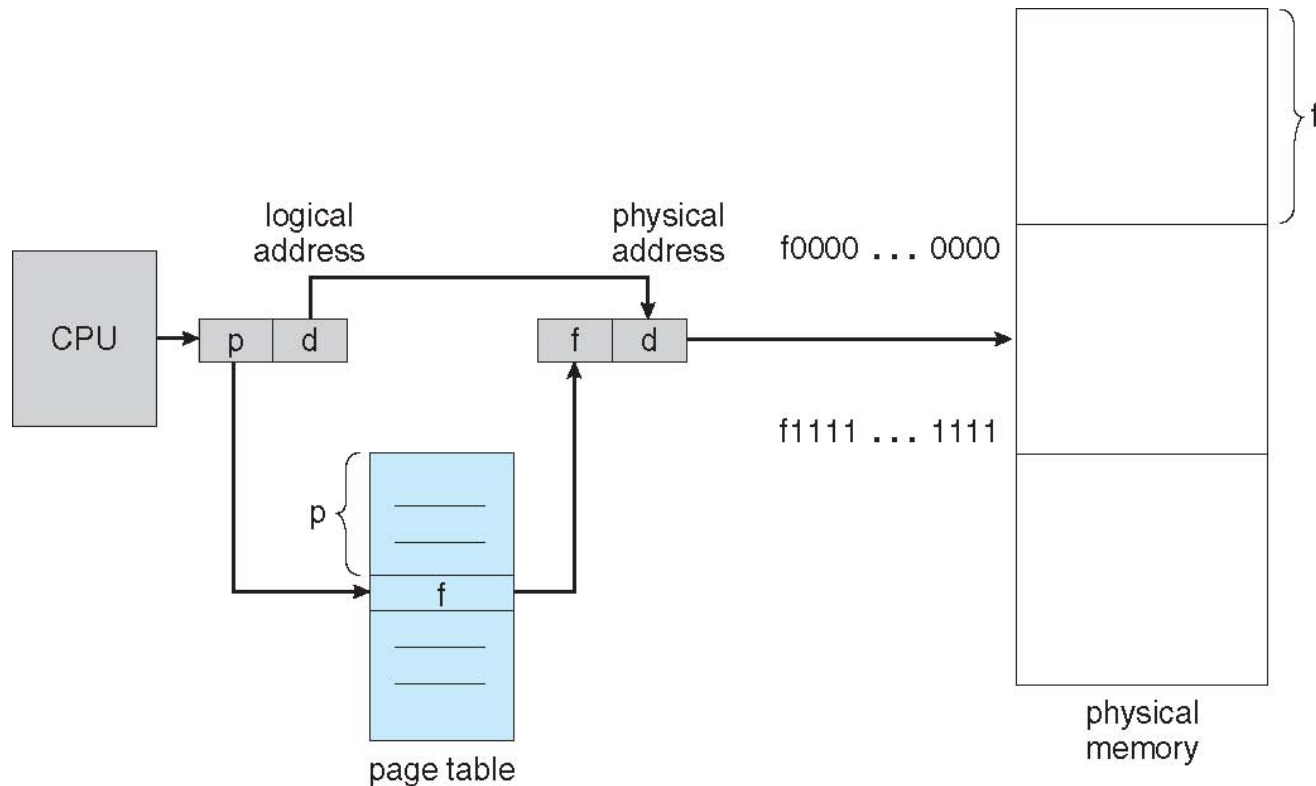
# Address Translation Scheme

- Address generated by CPU is divided into:
  - **Page number** ( $p$ ) – used as an index into a **page table** which contains base address of each page in physical memory
  - **Page offset** ( $d$ ) – combined with base address to define the physical memory address that is sent to the memory unit



- For given logical address space  $2^m$  and page size  $2^n$

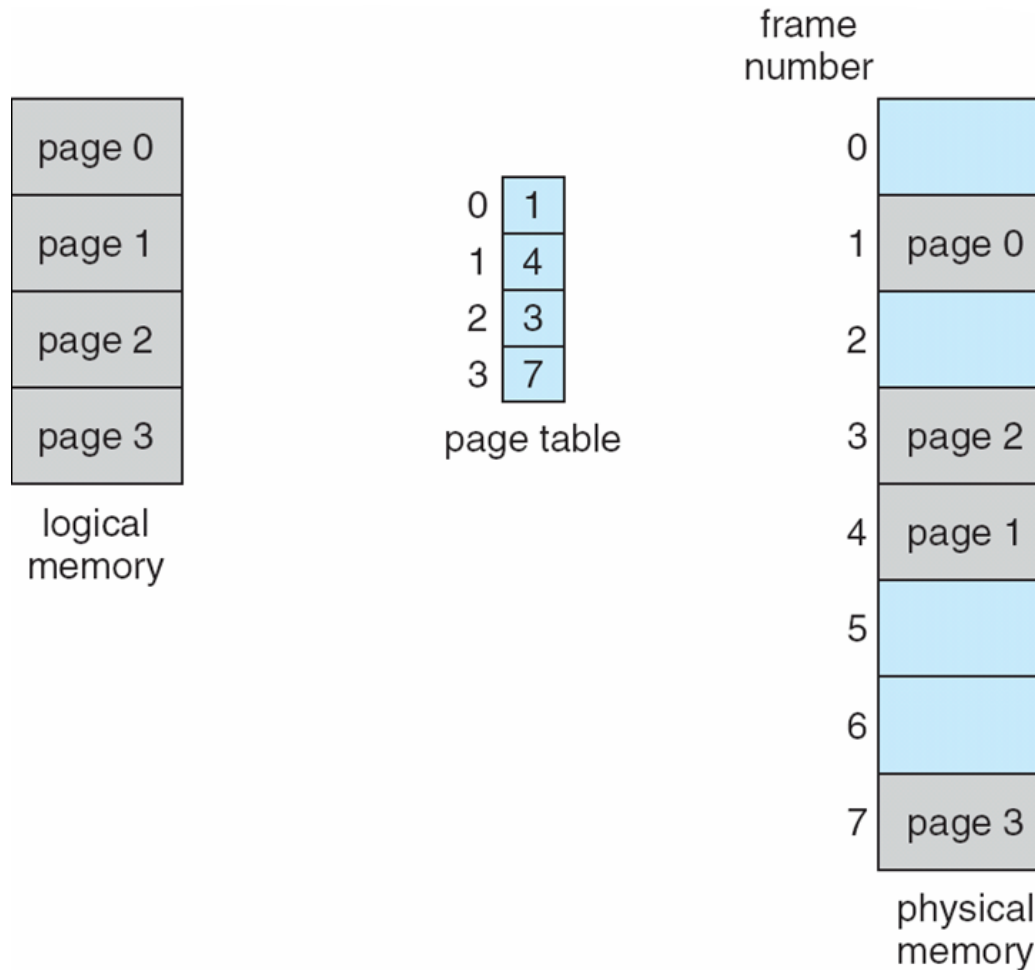
# Paging Hardware



Page number  $p$  to frame number  $f$  translation



# Paging Model of Logical and Physical Memory



# Paging Example

0	a
1	b
2	c
3	d
4	e
5	f
6	g
7	h
8	i
9	j
10	k
11	l
12	m
13	n
14	o
15	p

logical memory

0	5
1	6
2	1
3	2

page table

Page 0 maps  
to frame 5

0	
4	i j k l
8	m n o p
12	
16	
20	a b c d
24	e f g h
28	

physical memory


8 frames  
Frame number 0-to-7

$n=2$  and  $m=4$  Logical add. space =  $2^4$  bytes,  $2^2=4$ -byte pages  
32-byte memory with 8 frames

# Paging (Cont.)


- Internal fragmentation
  - Ex: Page size = 2,048 bytes, Process size = 72,766 bytes
    - 35 pages + 1,086 bytes
    - Internal fragmentation of  $2,048 - 1,086 = 962$  bytes
  - Worst case fragmentation = 1 frame – 1 byte
  - On average fragmentation =  $1 / 2$  frame size
  - So small frame sizes desirable?
    - But each page table entry takes memory to track
  - Page size growing over time
    - X86-64: 4 KB (common), 2 MB (“huge” for servers), 1GB (“large”)
- Process view and physical memory now very different
- By implementation, a process can only access its own memory unless ..

# Free Frame allocation



(a)

Before allocation



(b)

After allocation

# Implementation of Page Table

Page table is kept in main memory

- **Page-table base register (PTBR)** points to the page table
- **Page-table length register (PTLR)** indicates size of the page table
- In this scheme every data/instruction access requires **two memory accesses**
  - One for the page table and one for the data / instruction

One page-table  
For each process

The *two memory access problem* can be solved by the use of a special fast-lookup hardware cache called **associative memory** or **translation look-aside buffers (TLBs)**

TLB: cache for  
page Table

# Caching: Concept

- Widely used concept:
  - keep small subset of information likely to be needed in near future in a fast accessible place

## Examples:

- Cache Memory (“Cache”):  
Cache for Main memory
- Browser cache: for browser
- Disk cache
- Cache for Page Table: TLB

## Challenges:

- Is the information in cache? Where?
- Hit rate vs cache size

# Implementation of Page Table (Cont.)

- Some TLBs store **address-space identifiers (ASIDs)** in each TLB entry – uniquely identifies each process to provide address-space protection for that process
  - Otherwise need to flush at every context switch
- TLBs typically small (64 to 1,024 entries)
- On a TLB miss, value is loaded into the TLB for faster access next time
  - Replacement policies must be considered
  - Some entries can be **wired down** for permanent fast access

TLB: cache for  
page Table

# Associative Memory

- Associative memory – parallel search


Page #	Frame #

- Address translation (p, d)
  - If p is in associative register, get frame # out
  - Otherwise get frame # from page table in memory





# Paging Hardware With TLB



TLB Miss: page table access may be done using hardware or software

# Effective Access Time

- Associative Lookup =  $\varepsilon$  time unit
  - Can be  $< 10\%$  of memory access time
- **Hit ratio =  $\alpha$** 
  - Hit ratio – percentage of times that a page number is found in the associative registers; ratio related to number of associative registers
- **Effective Access Time (EAT):** probability weighted
$$\text{EAT} = (100 + \varepsilon) \alpha + (200 + \varepsilon)(1 - \alpha)$$
- Ex:

Consider  $\alpha = 80\%$ ,  $\varepsilon$  = negligible for TLB search, 100ns for memory access

  - $\text{EAT} = 0.80 \times 100 + 0.20 \times 200 = 120\text{ns}$
- Consider more realistic hit ratio  $\rightarrow \alpha = 99\%$ ,
  - $\text{EAT} = 0.99 \times 100 + 0.01 \times 200 = 101\text{ns}$

# Memory Protection

- Memory protection implemented by associating protection bit with each frame to indicate if read-only or read-write access is allowed
  - Can also add more bits to indicate page execute-only, and so on
- **Valid-invalid** bit attached to each entry in the page table:
  - “valid” indicates that the associated page is in the process’ logical address space, and is thus a legal page
  - “invalid” indicates that the page is not in the process’ logical address space
- Any violations result in a trap to the kernel

# Valid (v) or Invalid (i) Bit In A Page Table

00000	page 0
	page 1
	page 2
	page 3
	page 4
10,468	page 5
12,287	

frame number		valid-invalid bit
0	2	v
1	3	v
2	4	v
3	7	v
4	8	v
5	9	v
6	0	i
7	0	i

page table

“invalid” : page is not in the process’s address space.

0	
1	
2	page 0
3	page 1
4	page 2
5	
6	
7	page 3
8	page 4
9	page 5
	⋮
	page $n$

# Shared Pages

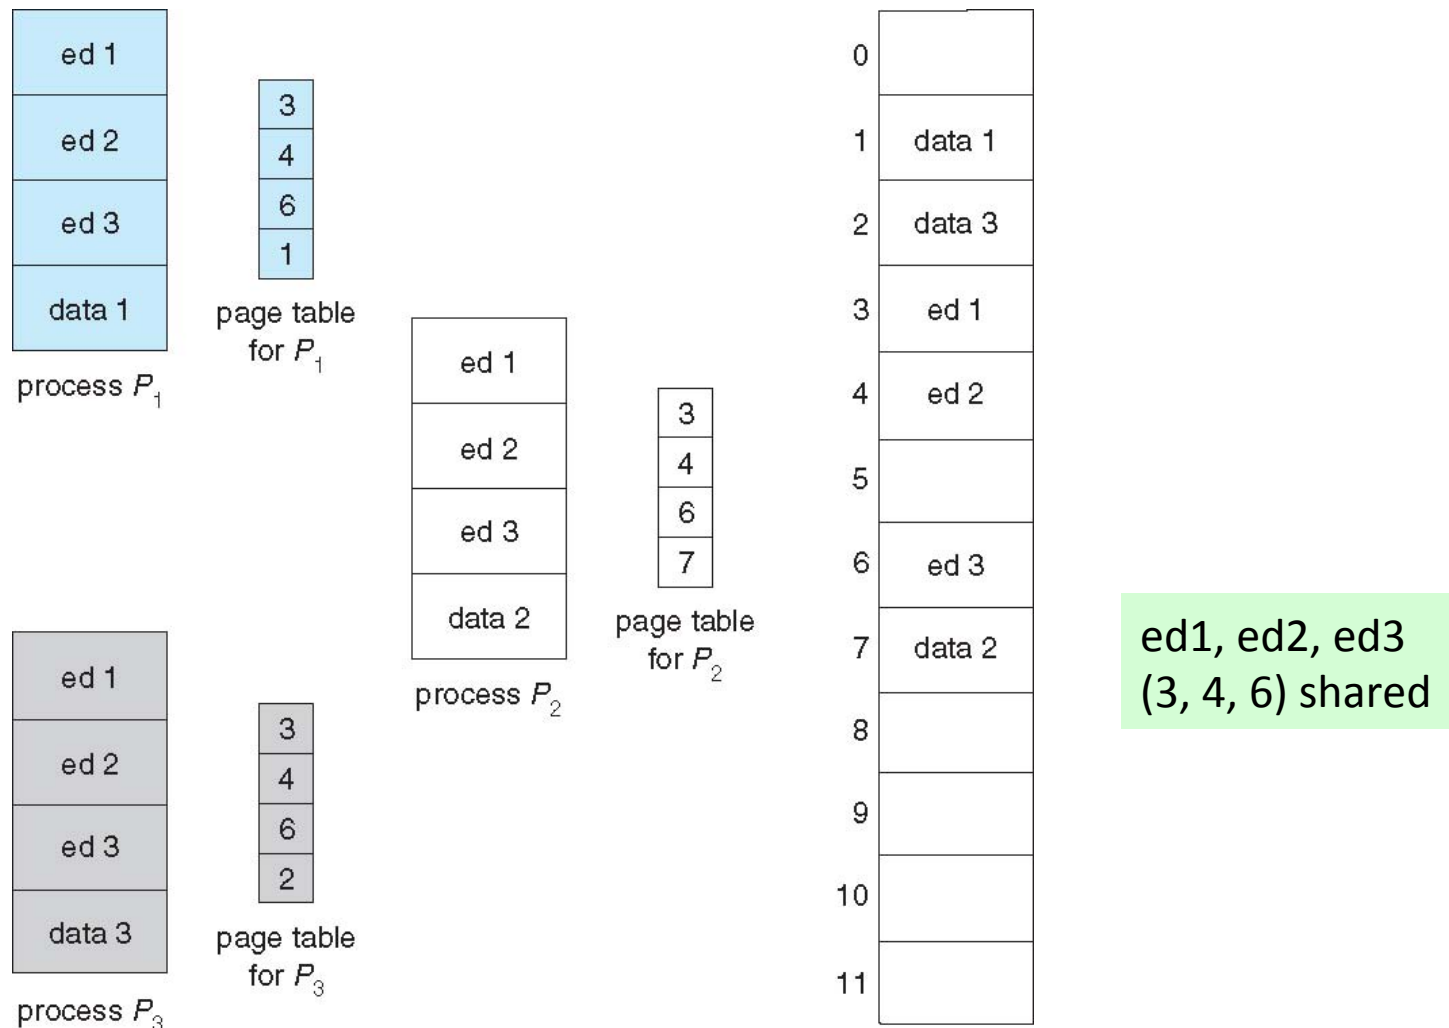
- **Shared code**

- One copy of read-only (**reentrant** non-self modifying) code *shared* among processes (i.e., text editors, compilers, window systems)
- Similar to multiple threads sharing the same process space
- Also useful for interprocess communication if sharing of read-write pages is allowed

- **Private code and data**

- Each process keeps a separate copy of the code and data
- The pages for the private code and data can appear anywhere in the logical address space

# Shared Pages Example



## Overheads in paging: Page table and internal fragmentation

### Optimal Page Size: page table size vs int fragmentation tradeoff

- Average process size =  $s$
- Page size =  $p$
- Size of each page entry =  $e$ 
  - Pages per process =  $s/p$
  - $se/p$ : Total page table space
- Total Overhead = Page table overhead + Internal fragmentation loss  
 $= se/p + p/2$

## Optimal Page size: Page table and internal fragmentation

- Total Overhead =  $se/p + p/2$
- Optimal: First derivative with respect to  $p$ , equate to 0

$$-se/p^2 + 1/2 = 0$$

- i.e.  $p^2 = 2se$  or  $p = (2se)^{0.5}$

**Assume**  $s = 128\text{KB}$  and  $e=8$  bytes per entry

- Optimal page size = 1448 bytes
  - In practice we will never use 1448 bytes
  - Instead, either 1K or 2K would be used
    - **Why?** Pages sizes are in powers of 2 i.e.  $2^x$
    - Deriving offsets and page numbers is also easier



# Page Table Size

- Memory structures for paging can get huge using straight-forward methods
  - Consider a 32-bit logical address space as on modern computers
  - Page size of 4 KB ( $2^{12}$ ) entries
  - Page table would have 1 million entries ( $2^{32} / 2^{12}$ )
  - If each entry is 4 bytes -> 4 MB of physical address space / memory for page table alone
    - That amount of memory used to cost a lot
    - Don't want to allocate that contiguously in main memory

$2^{10}$	1024 or 1 kibibyte
$2^{20}$	1M mebibyte
$2^{30}$	1G gigibyte

# Issues with large page tables

- Cannot allocate page table **contiguously** in memory
- Solutions:
  - Divide the page table into smaller pieces
  - **Page the page-table**
    - Hierarchical Paging

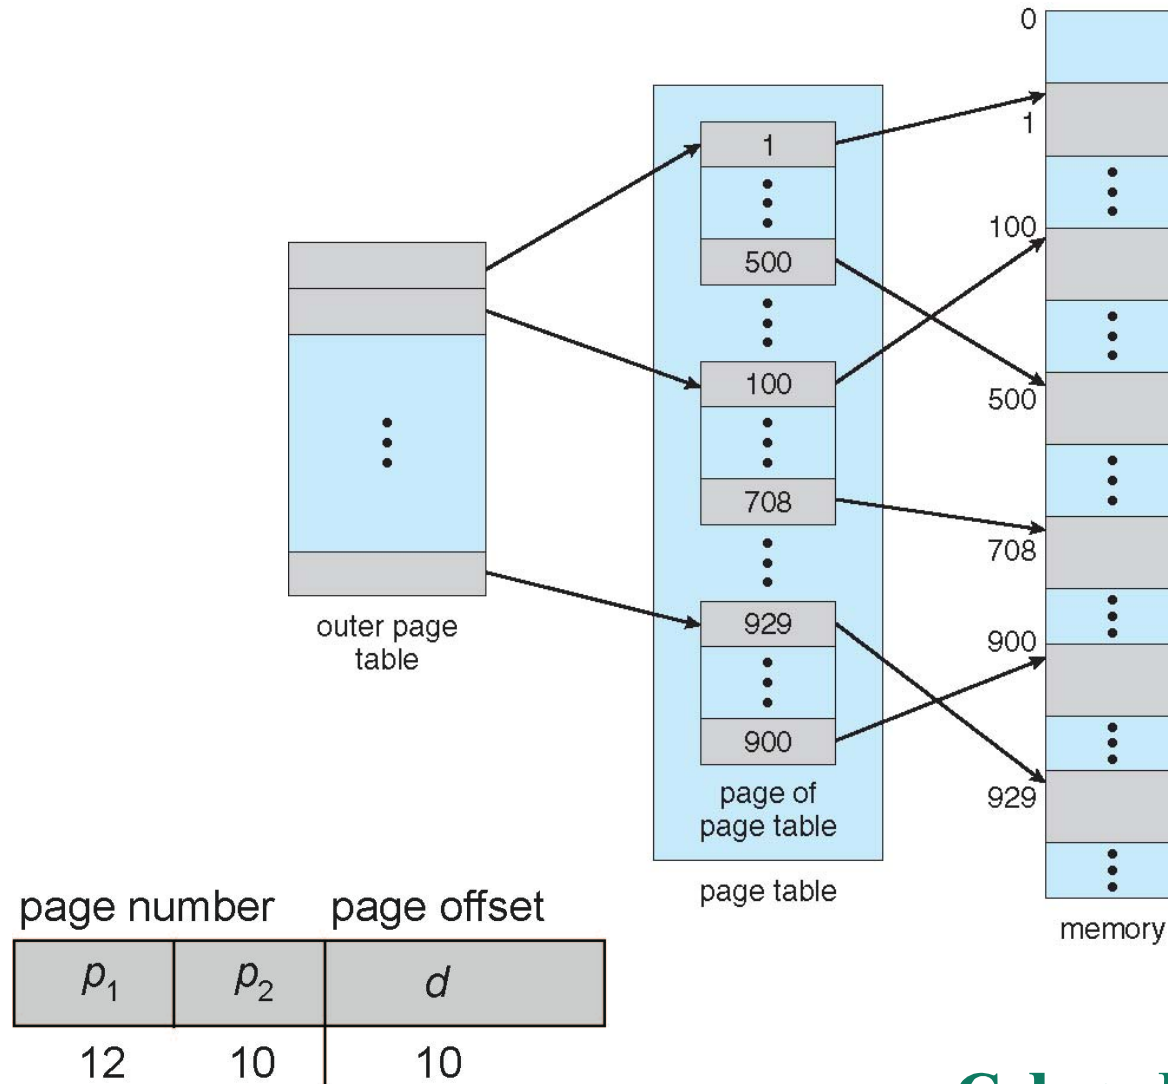
# Hierarchical Page Tables

- Break up the logical address space into multiple page tables
- A simple technique is a two-level page table
- We then page the page table

page number		page offset
$p_1$	$p_2$	$d$
12	10	10


P1: indexes the outer page table  
P2: page table: maps to frame

# Two-Level Page-Table Scheme



# Two-Level Paging Example

- A logical address (on 32-bit machine with 1K page size) is divided into:
  - a page number consisting of 22 bits
  - a page offset consisting of 10 bits
- Since the page table is paged, the page number is further divided into:
  - a 12-bit page number
  - a 10-bit page offset
- Thus, a logical address is as follows:



- where  $p_1$  is an index into the outer page table, and  $p_2$  is the displacement within the page of the inner page table
- Known as **forward-mapped page table**